IMEN891K / IDSC501 / AXGS-501 · DATA ENGINEERING · FALL 2026

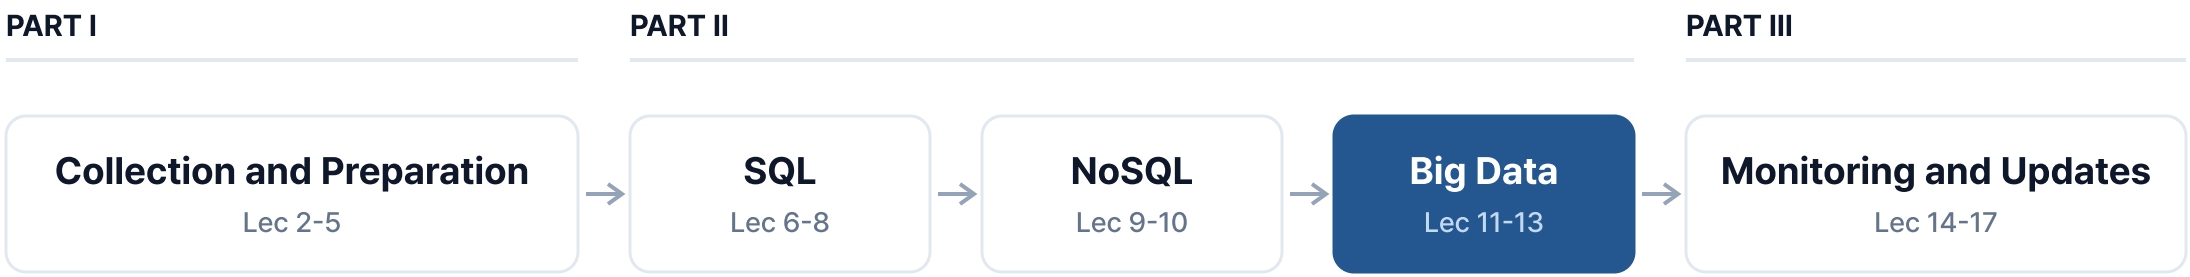
13.

Big Data III: Stream Processing

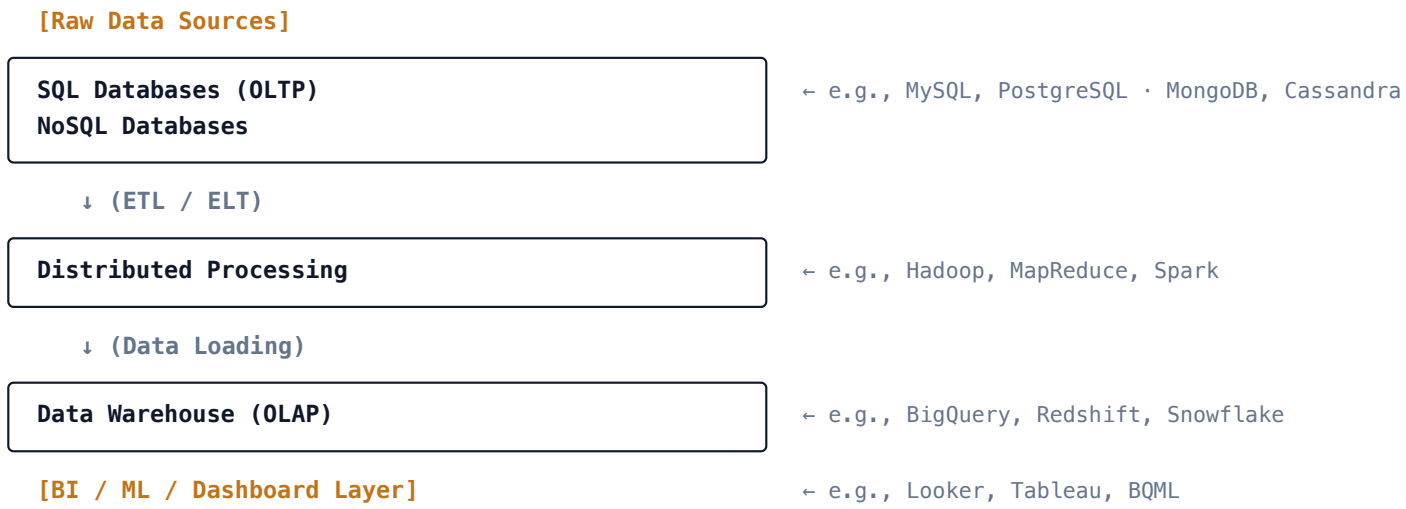
Sangwoo Mo · POSTECH

Industrial & Management Engineering

Where We Are


PART I
PART II
PART III
Collection and Preparation
Lec 2-5
SQL
Lec 6-8
NoSQL
Lec 9-10
Big Data
Lec 11-13
Monitoring and Updates
Lec 14-17

Recap: Where the data has been so far



Databases (OLTP) → processing (MapReduce, Spark) → warehouse or lakehouse (OLAP) → dashboards, with Airflow scheduling the steps. Every arrow so far moved data **in batches**.

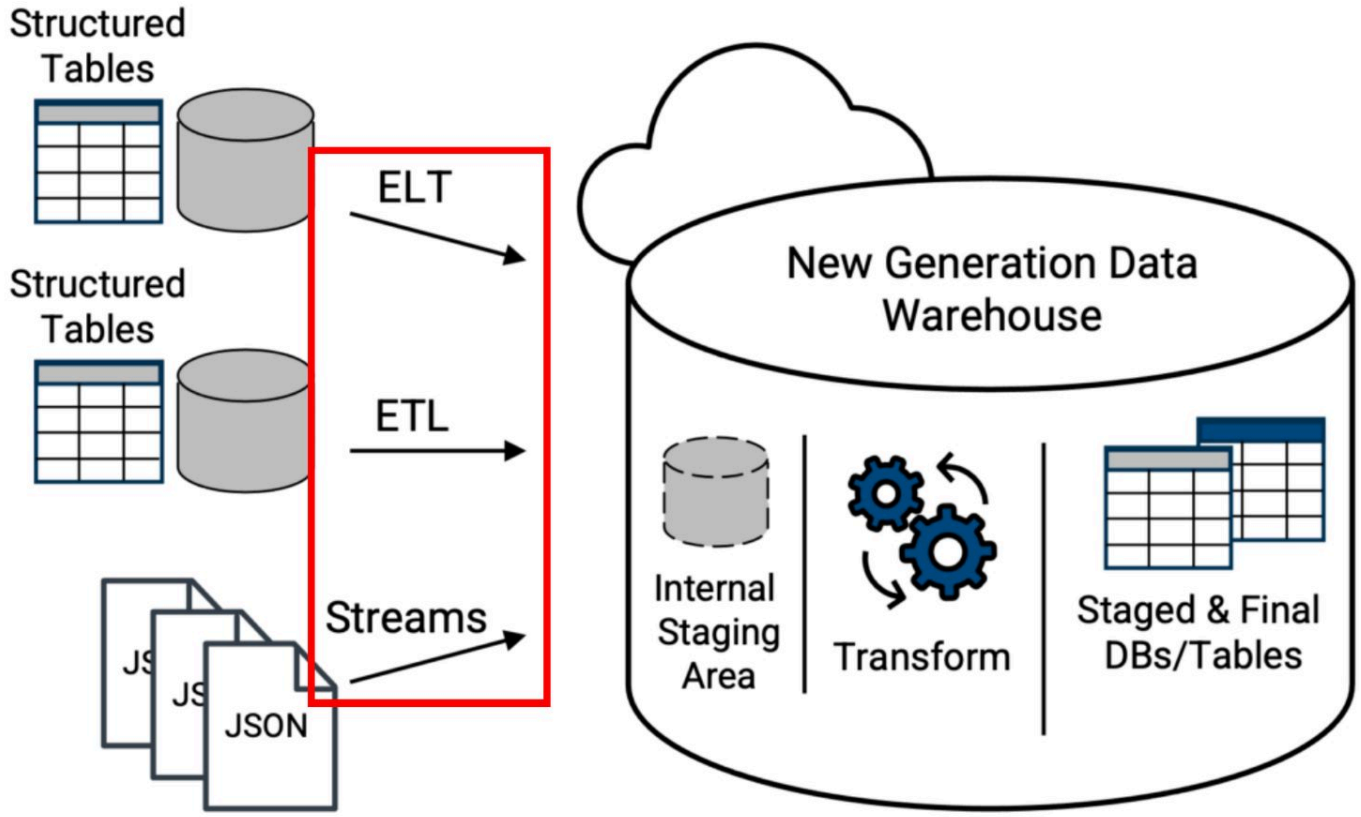
Outline: Big Data III

01	**From Batch to Live Events**	handle events as they arrive
02	**Moving Events with Kafka**	store events and deliver them reliably
03	**Processing Event Streams**	connect sources and process live data

01 · From Batch to Live Events

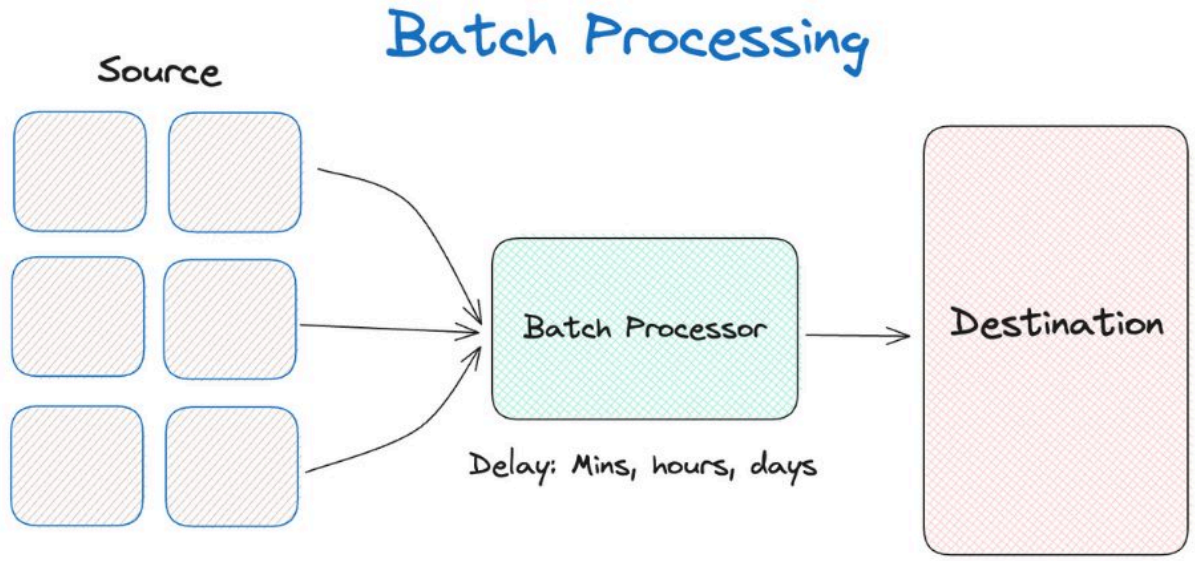
01	**From Batch to Live Events**	handle events as they arrive
02	**Moving Events with Kafka**	store events and deliver them reliably
03	**Processing Event Streams**	connect sources and process live data

Streams in the Data Flow



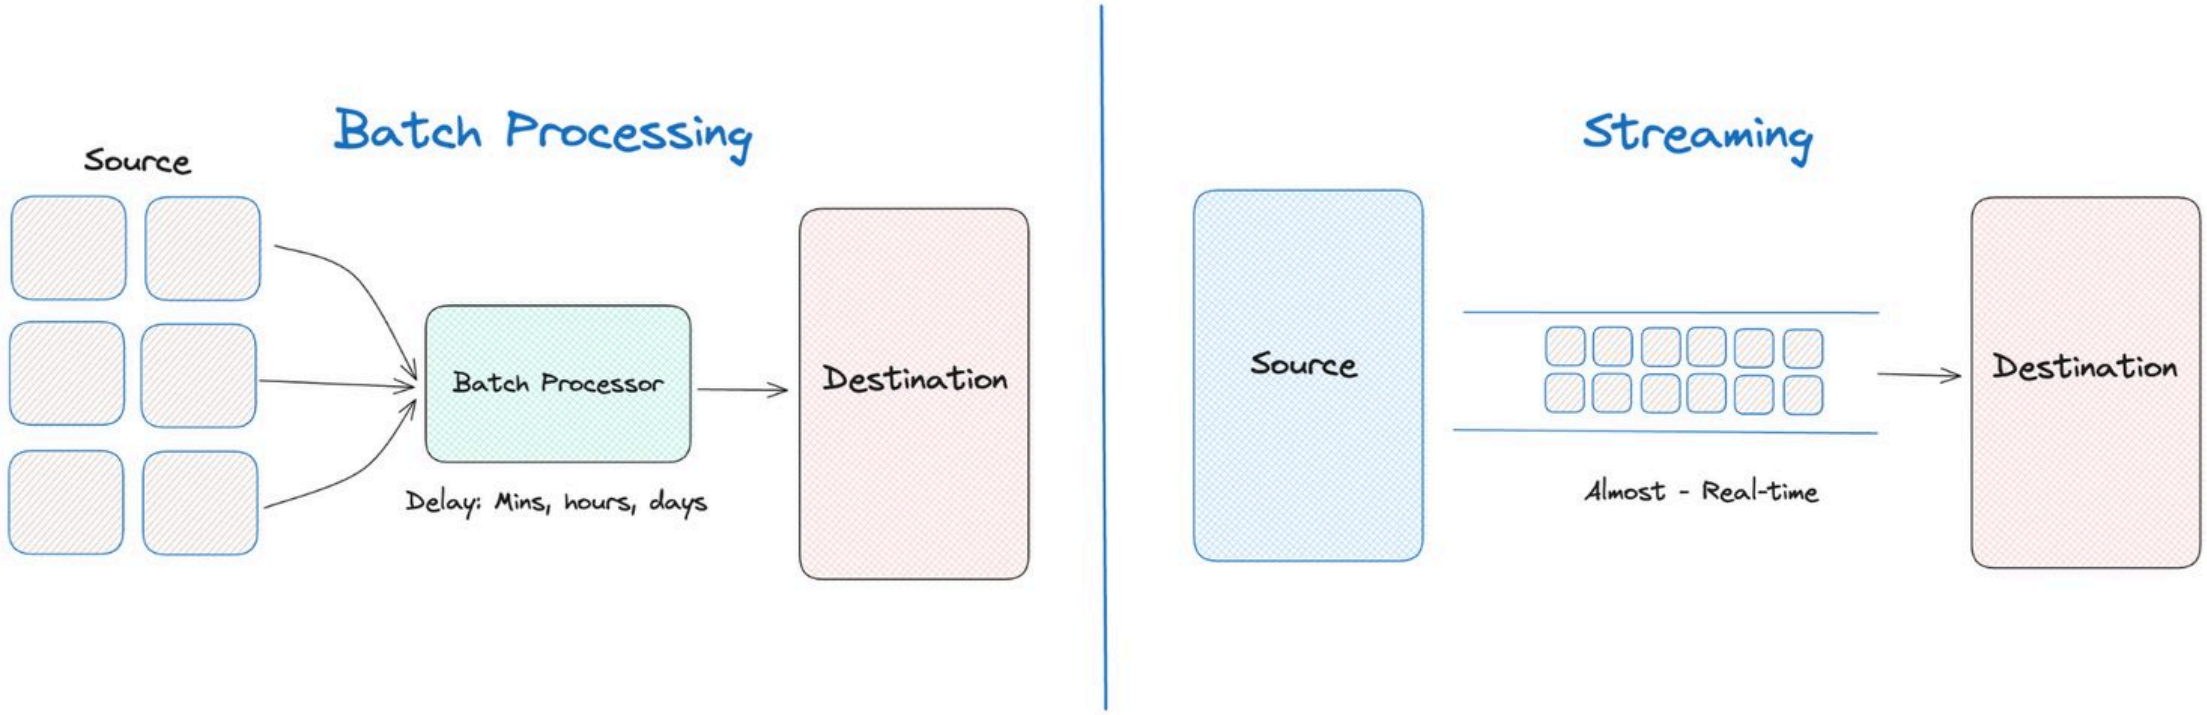
Three arrows in; the third, **Streams**, is today

Batch Processing



Batch: collect, then process, then deliver. **Delay: minutes to days**

Stream Processing



Streaming: each record flows as it happens. **Almost real-time**

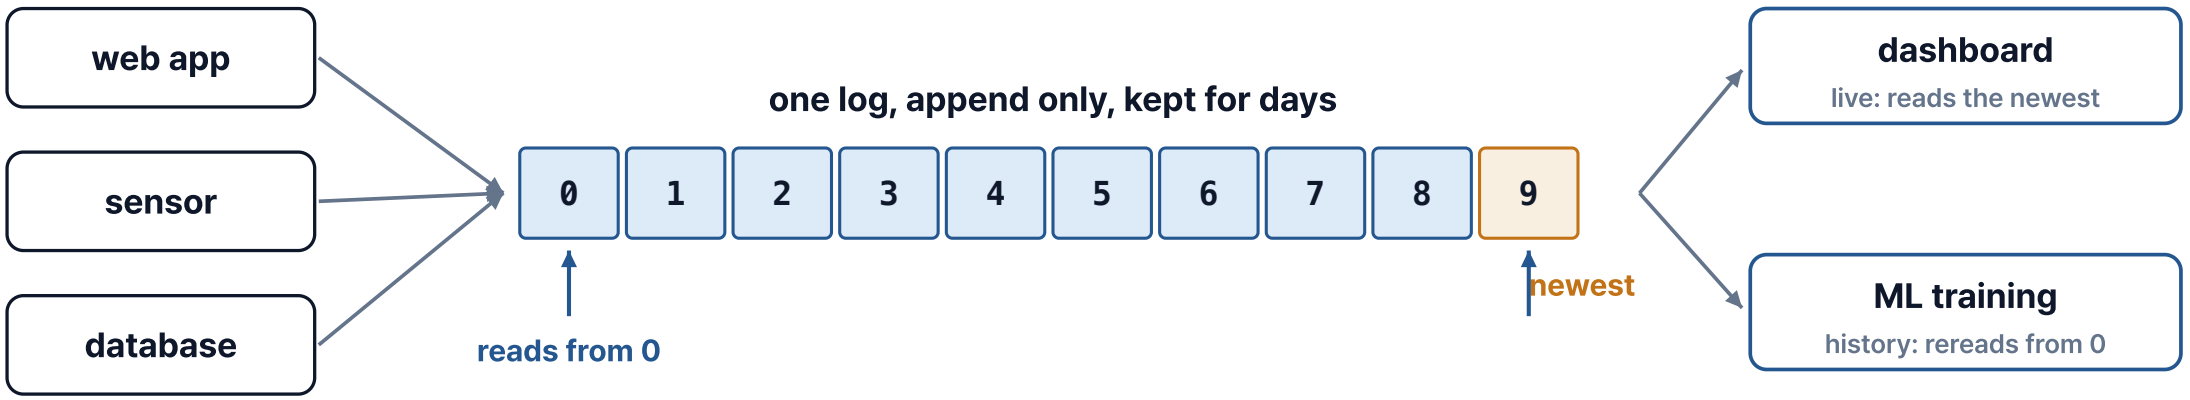
From Snapshots to Events



From **a snapshot now and then** to **a stream of every event**

Why an Event Log

One platform for every event, delivered live and kept for replay



Every system writes once, every reader picks its own position: **live and history from one place**

Readers choose their own position: **the newest event or the first**

Case study: Coupang Eats Data Flow

coupang eats data ingestion and pre-processing

Flow	Ingestion	Processing	Data engine
Non-real-time	ETL software · Kafka · cloud storage	Hive, job scheduler	Offline data engine (velocity / relationship)
Near real-time	Kafka	OLAP engine	NRT data engine (velocity)
Pure real-time	Kafka · feature service API	Spark / distributed processing · real-time consumer · Kafka	Pure real-time data engine (velocity / relationship / LBS)

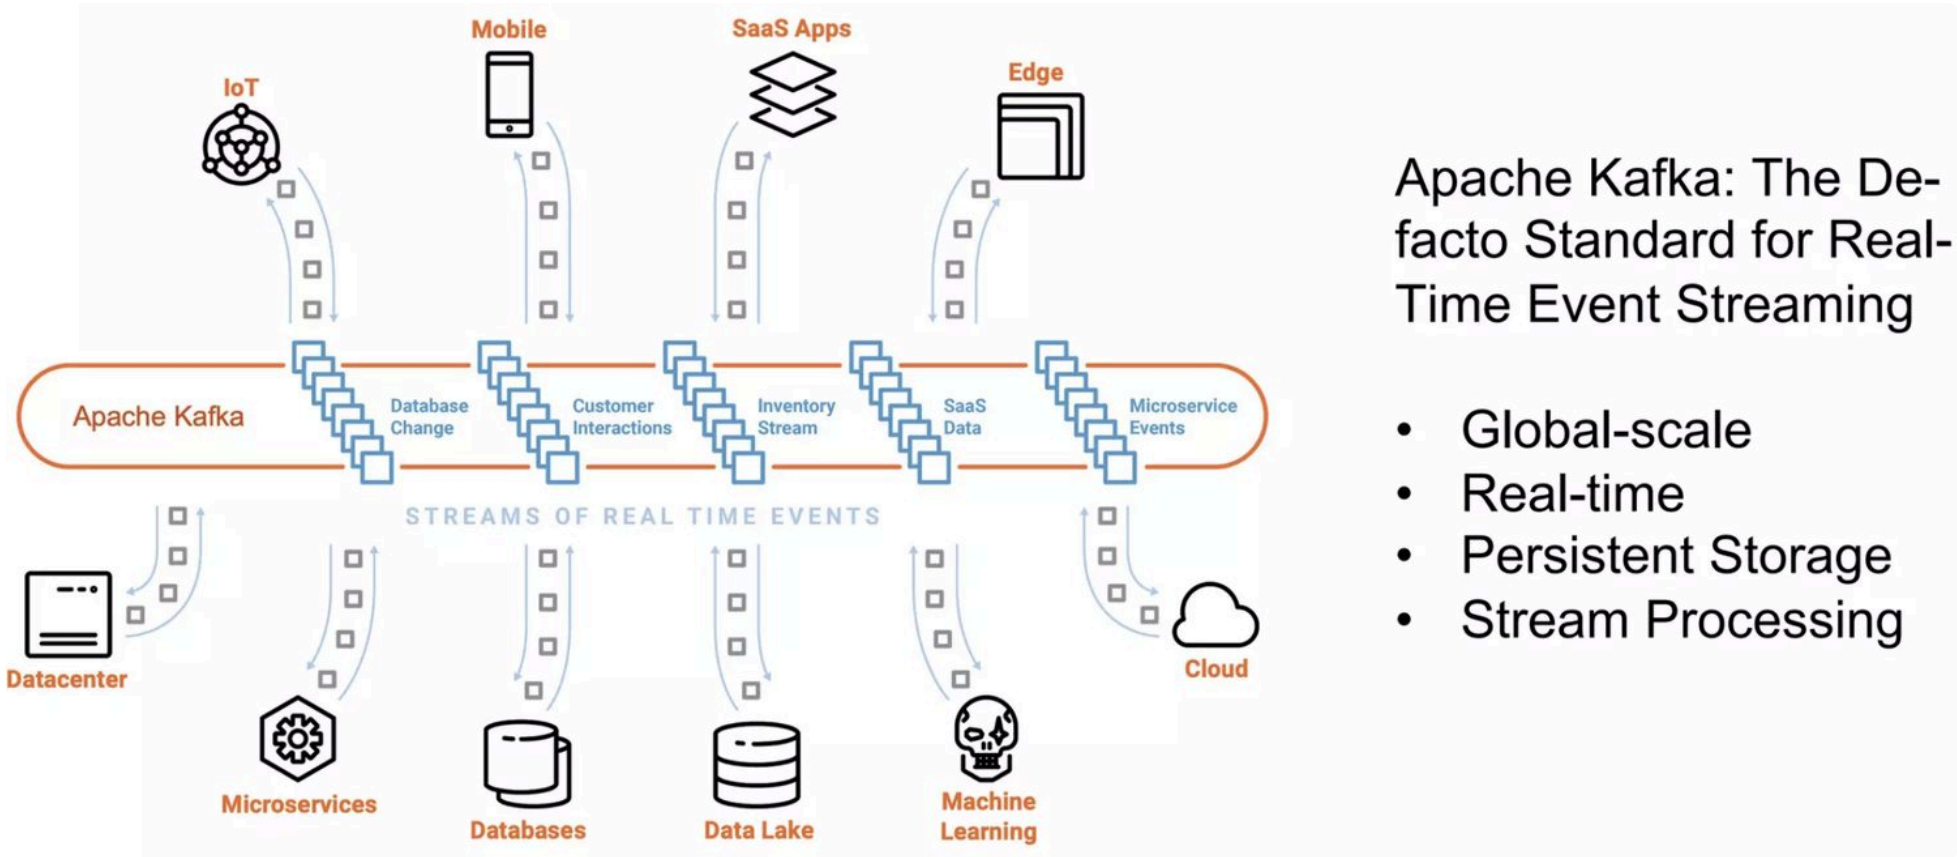
Three latencies, one ingestion bus: **Kafka** appears in every row.

Kafka carries events into batch, near-real-time, and live processing paths

02 · Moving Events with Kafka

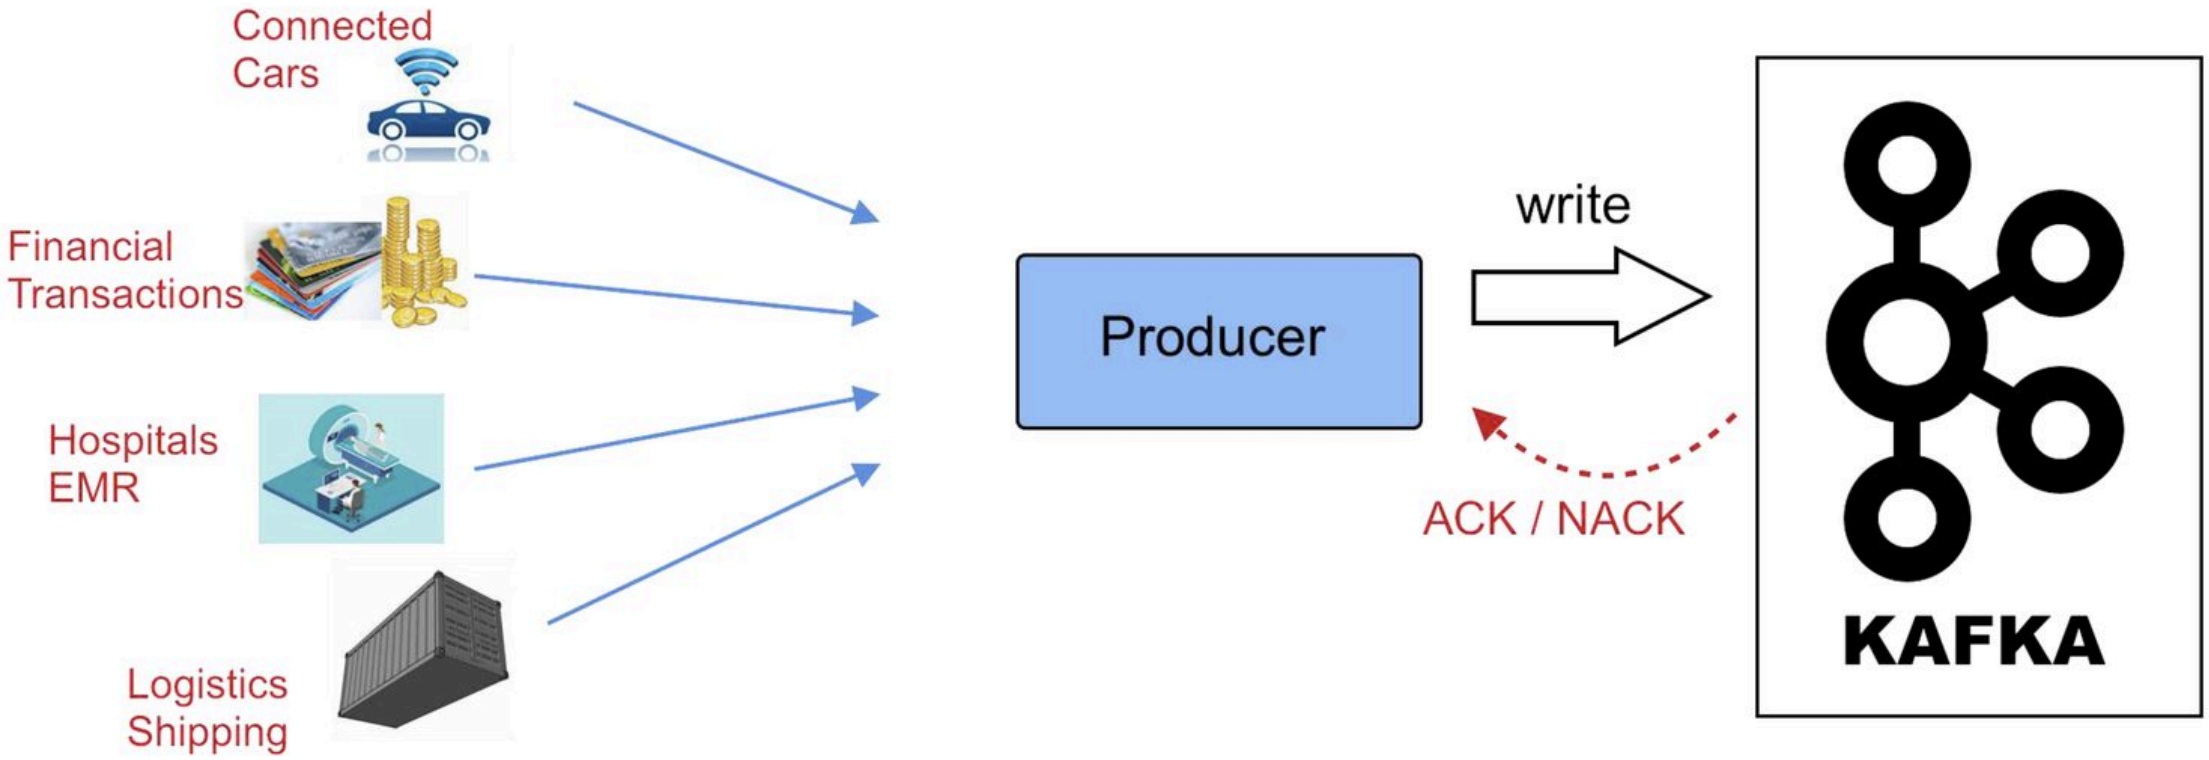
01	From Batch to Live Events	handle events as they arrive
02	**Moving Events with Kafka**	store events and deliver them reliably
03	Processing Event Streams	connect sources and process live data

Kafka



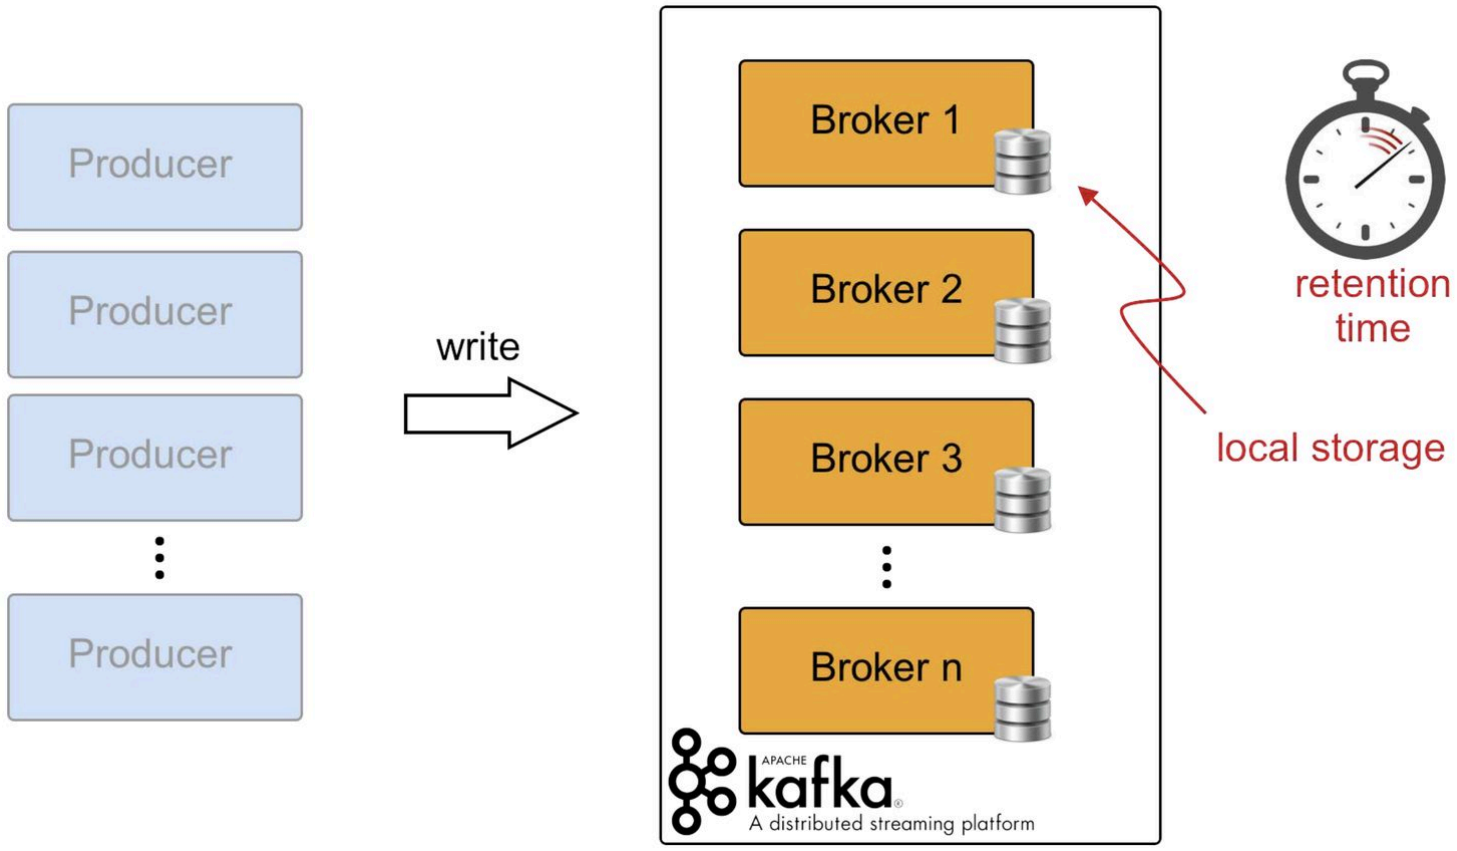
Everything that emits events on top, everything that wants them below, **one log in between**

Kafka – Producers



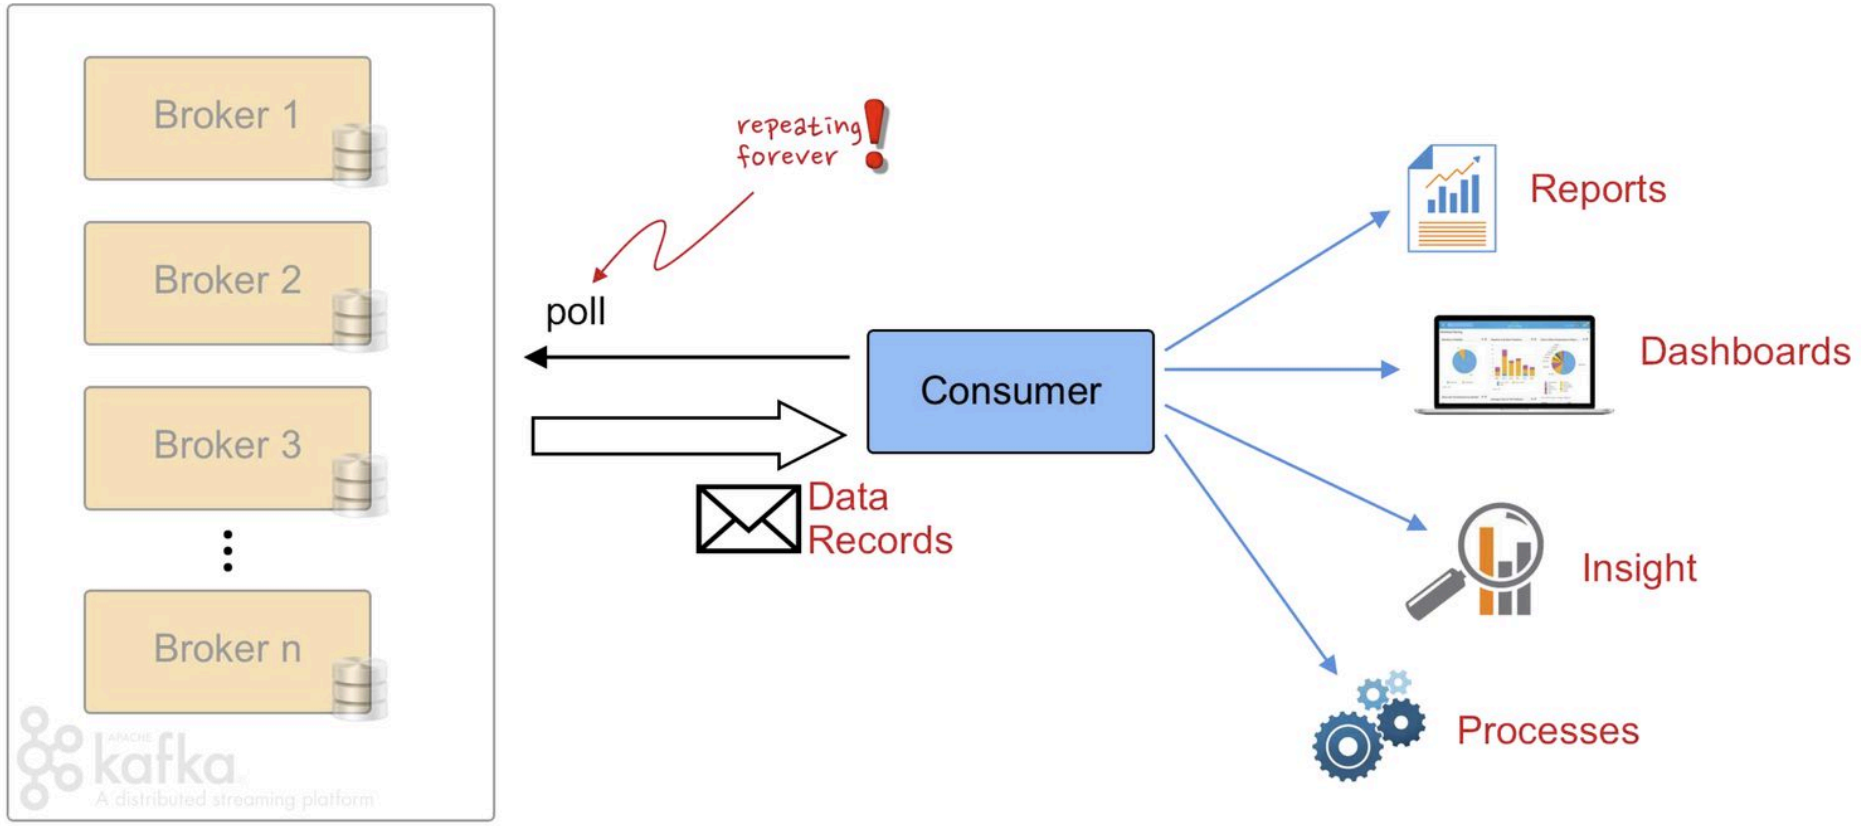
Producers write and get an acknowledgement; **they never know the readers**

Kafka – Brokers



Brokers store events on local disk for a **retention time**: a store, not a queue

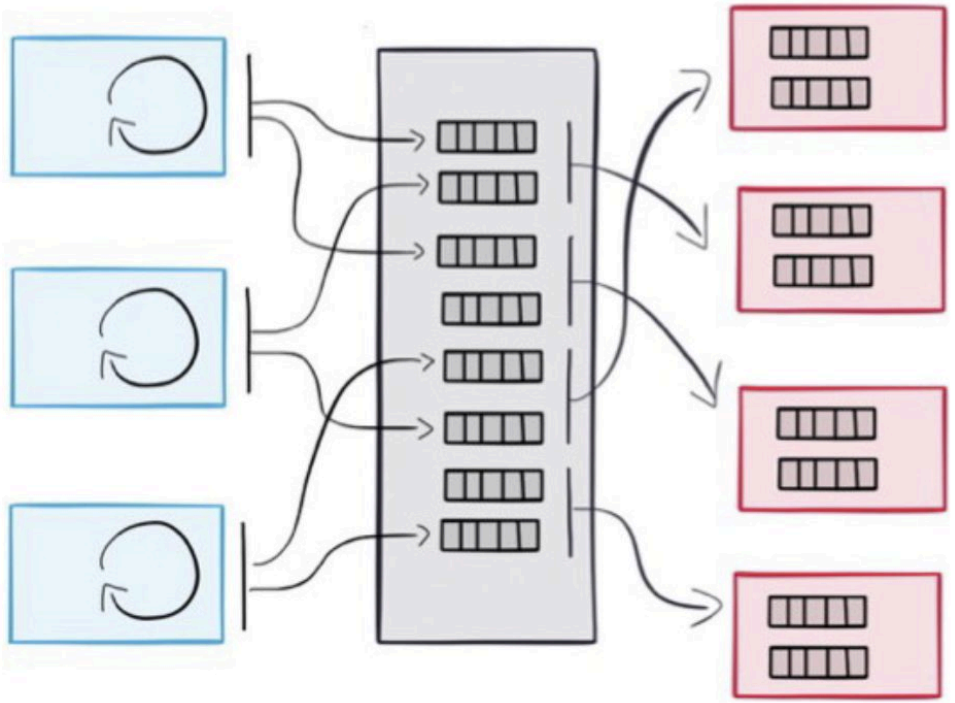
Kafka – Consumers



Consumers **poll**, forever, at their own pace

Decoupling Producers and Consumers

- Producers and Consumers are decoupled
- Slow Consumers do not affect Producers
- Add Consumers without affecting Producers
- Failure of Consumer does not affect System

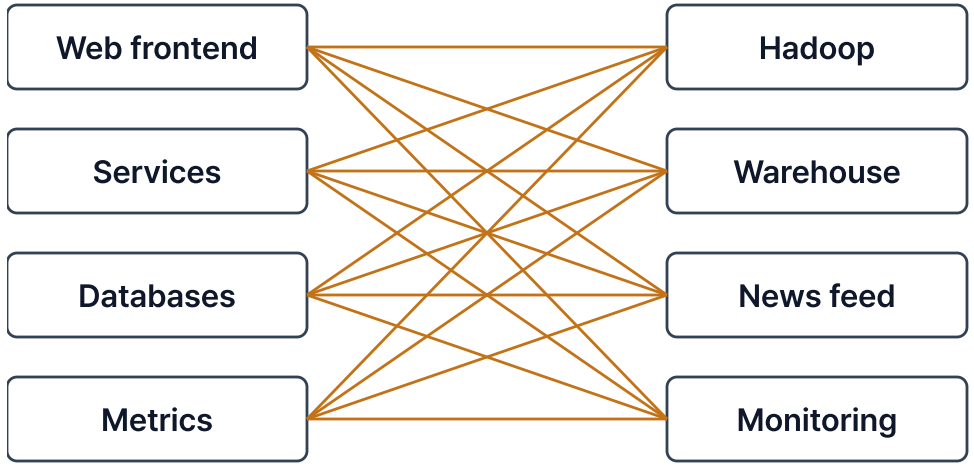


The log in the middle lets each side **fail, slow down or grow** alone

Case study: Kafka at LinkedIn

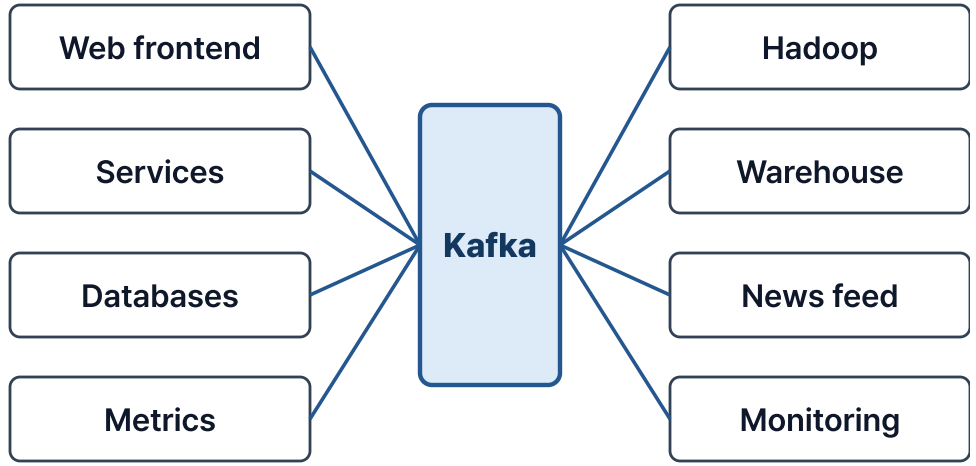
LinkedIn open-sourced Kafka in 2011 to distribute activity and operational events

Illustrative: All-to-All Pipelines



4 × 4 = 16 custom pipelines

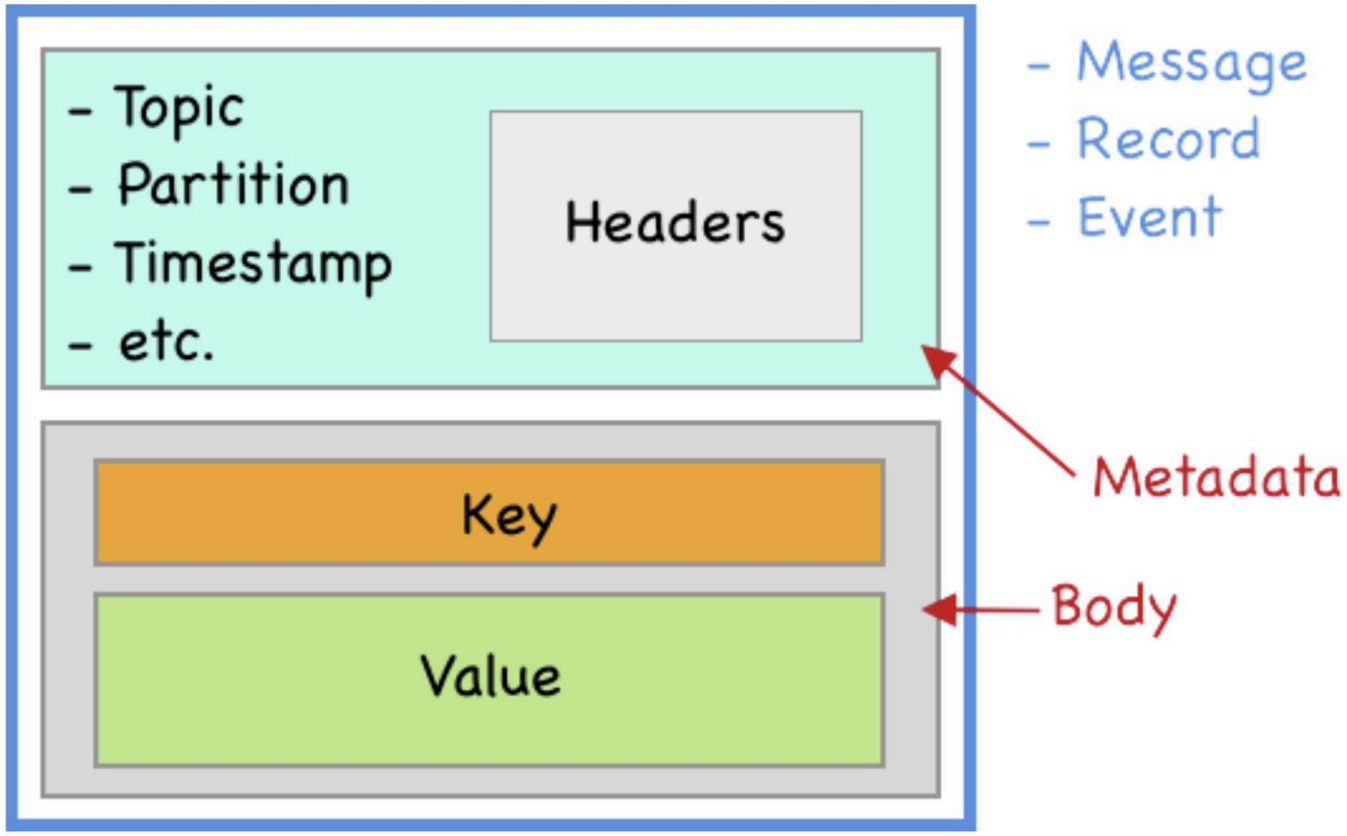
After: One Shared Log



4 + 4 = 8 connections to one log

With all-to-all integration, a shared log reduces endpoint connections from **n×m** to **n+m**

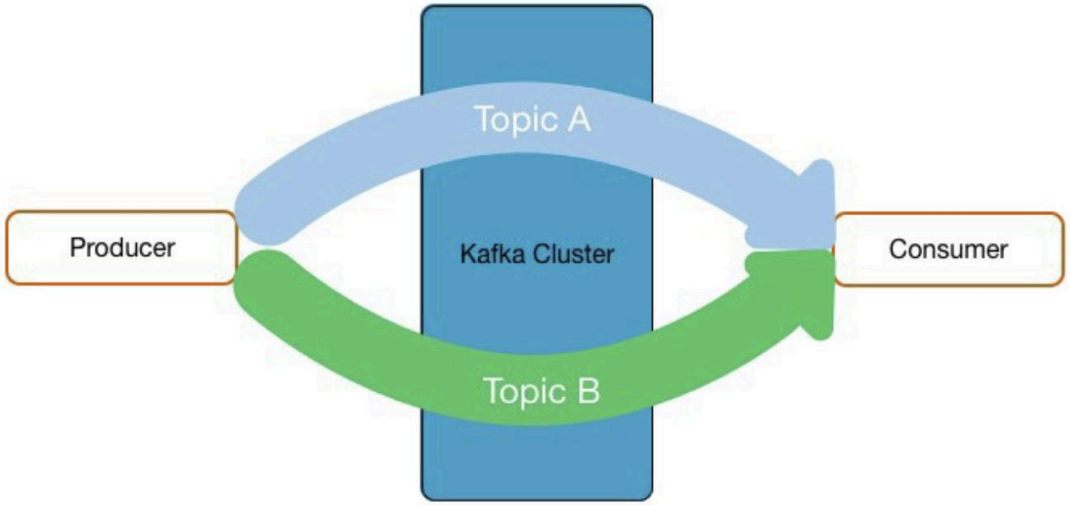
Data Elements



Metadata and optional headers on top; **key and value**, both bytes, as the body

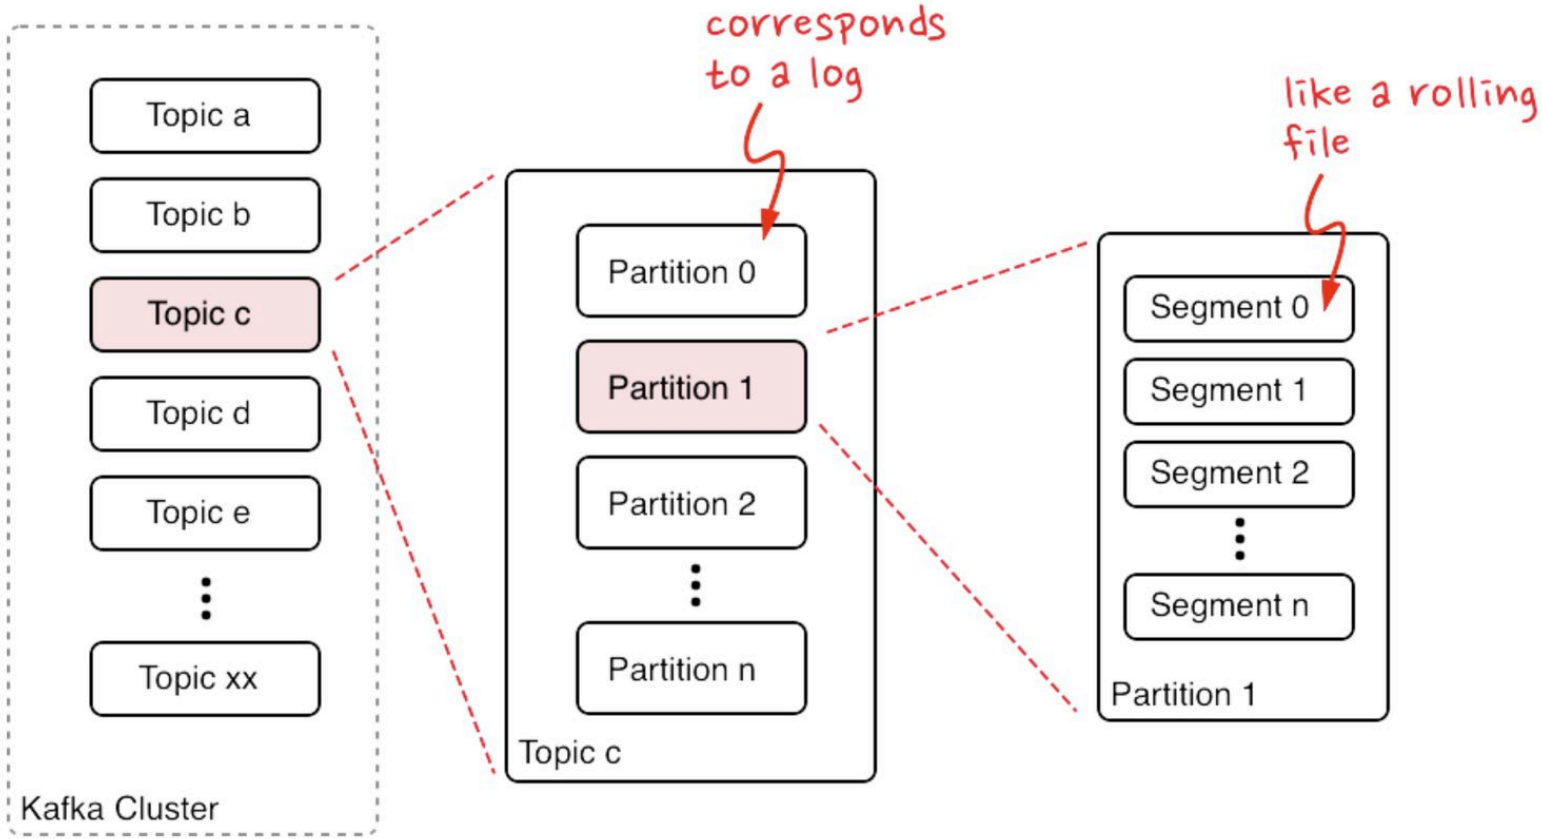
Topics

- **Topics**: streams of “related” messages in Kafka
  - Is a logical representation
  - Categorizes messages into groups
- Developers define topics
- Producer ↔ topic: N to N relation
- Unlimited number of topics



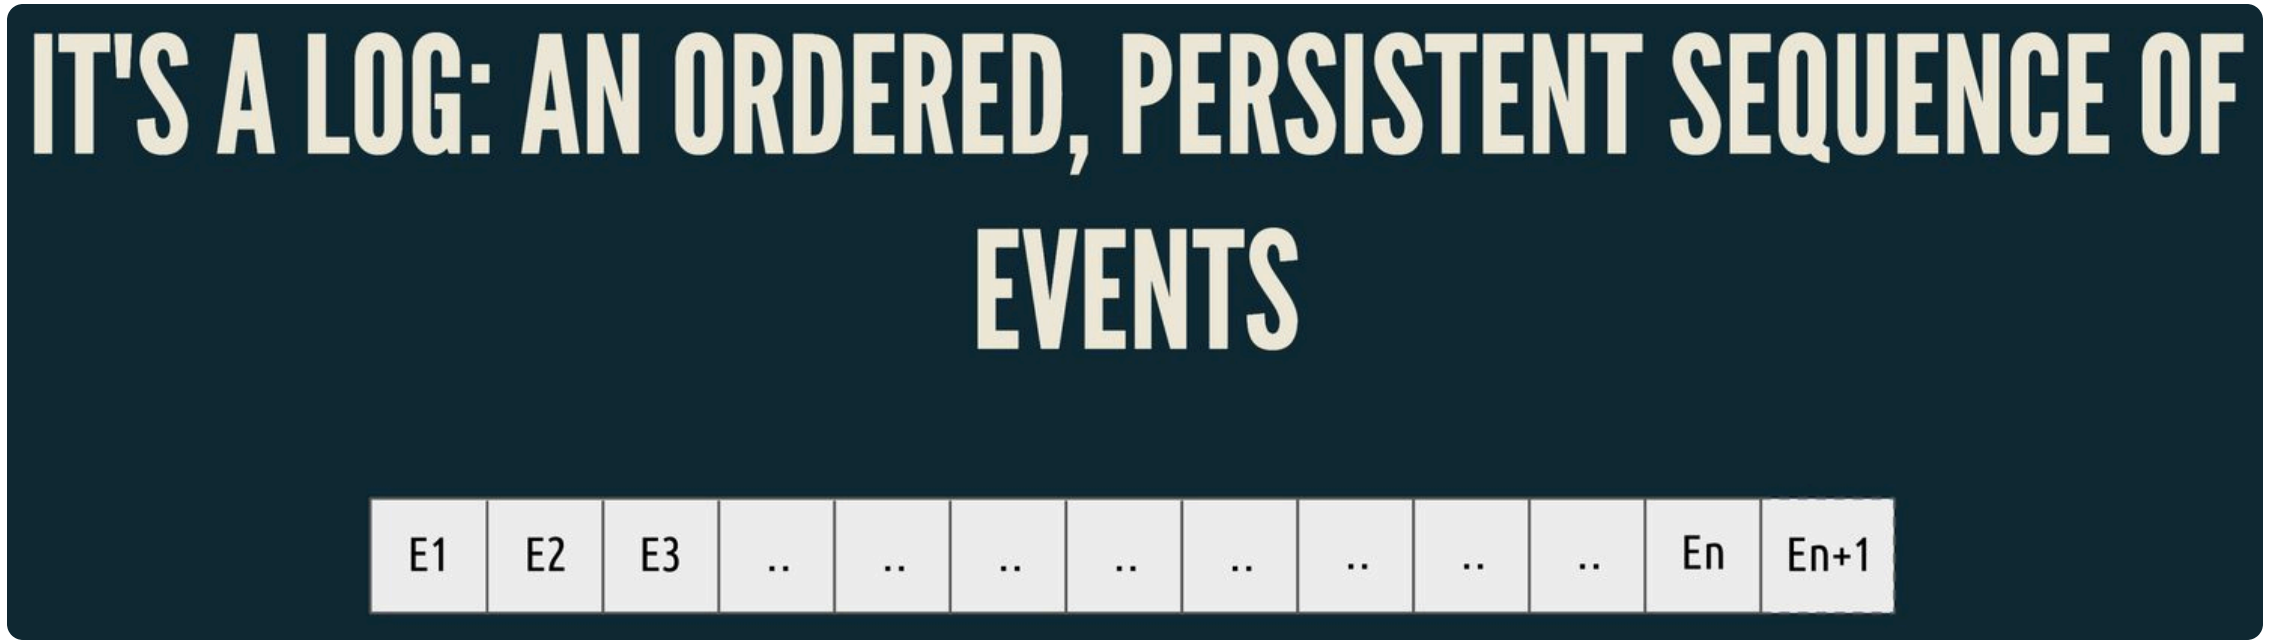
A topic is a **named stream**: the nearest thing to a table

Topic, Partition, Segment



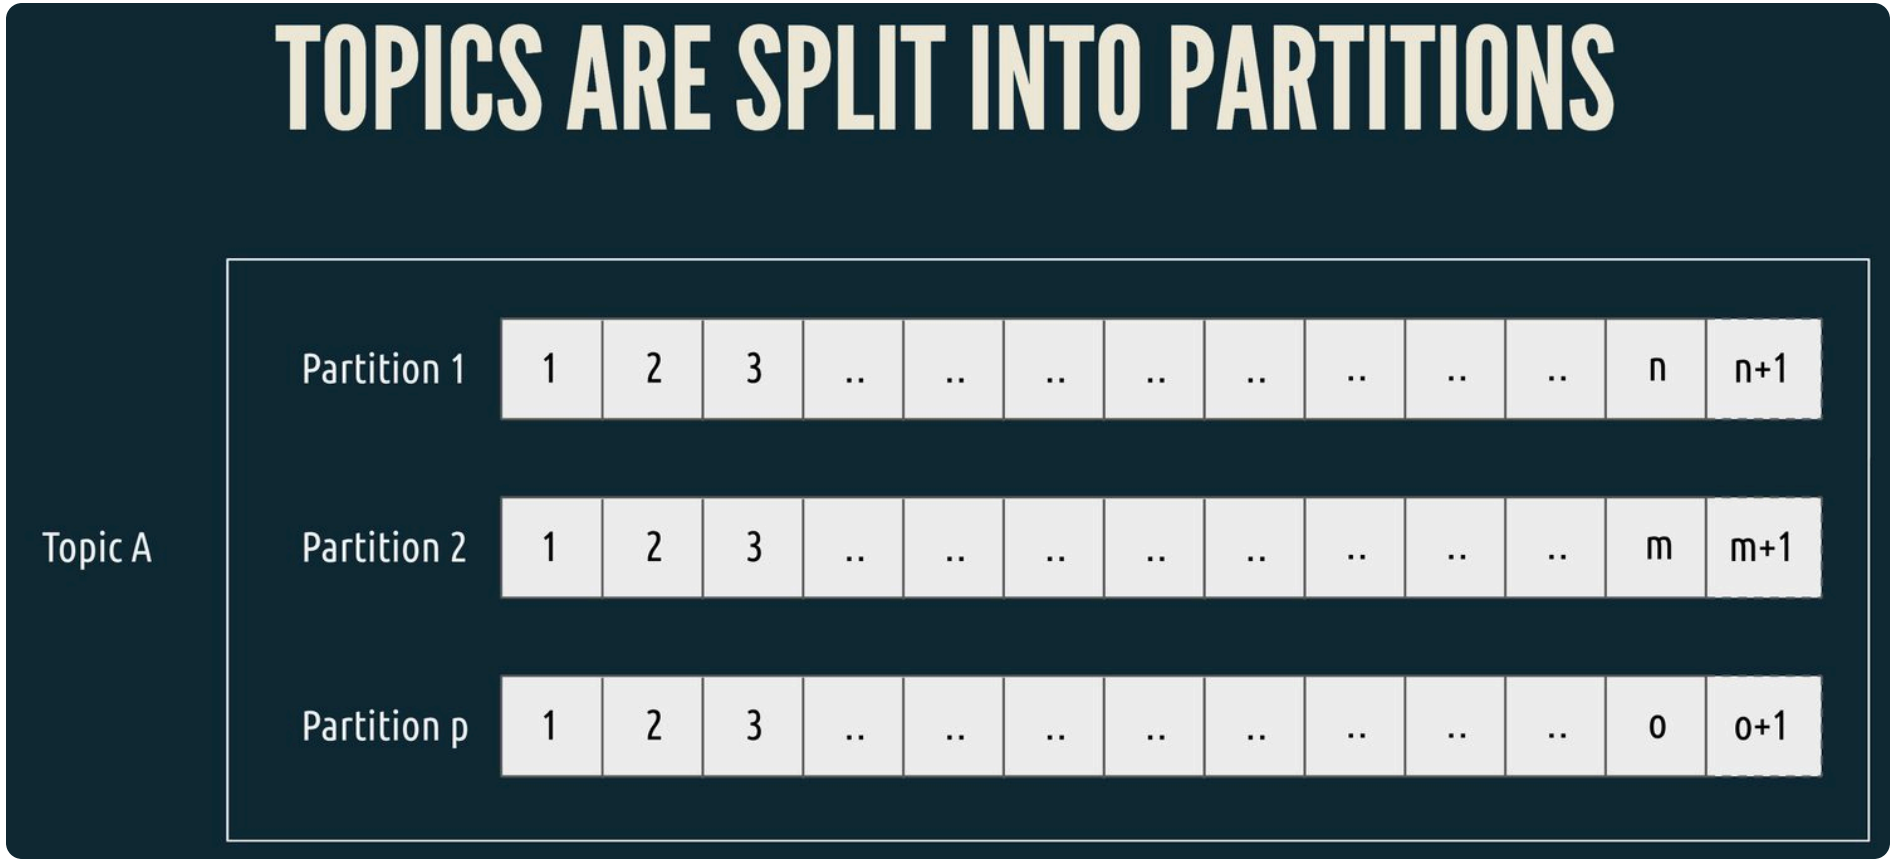
Topic → partitions (each a log) → segments (rolling files on disk)

An Append-Only Log



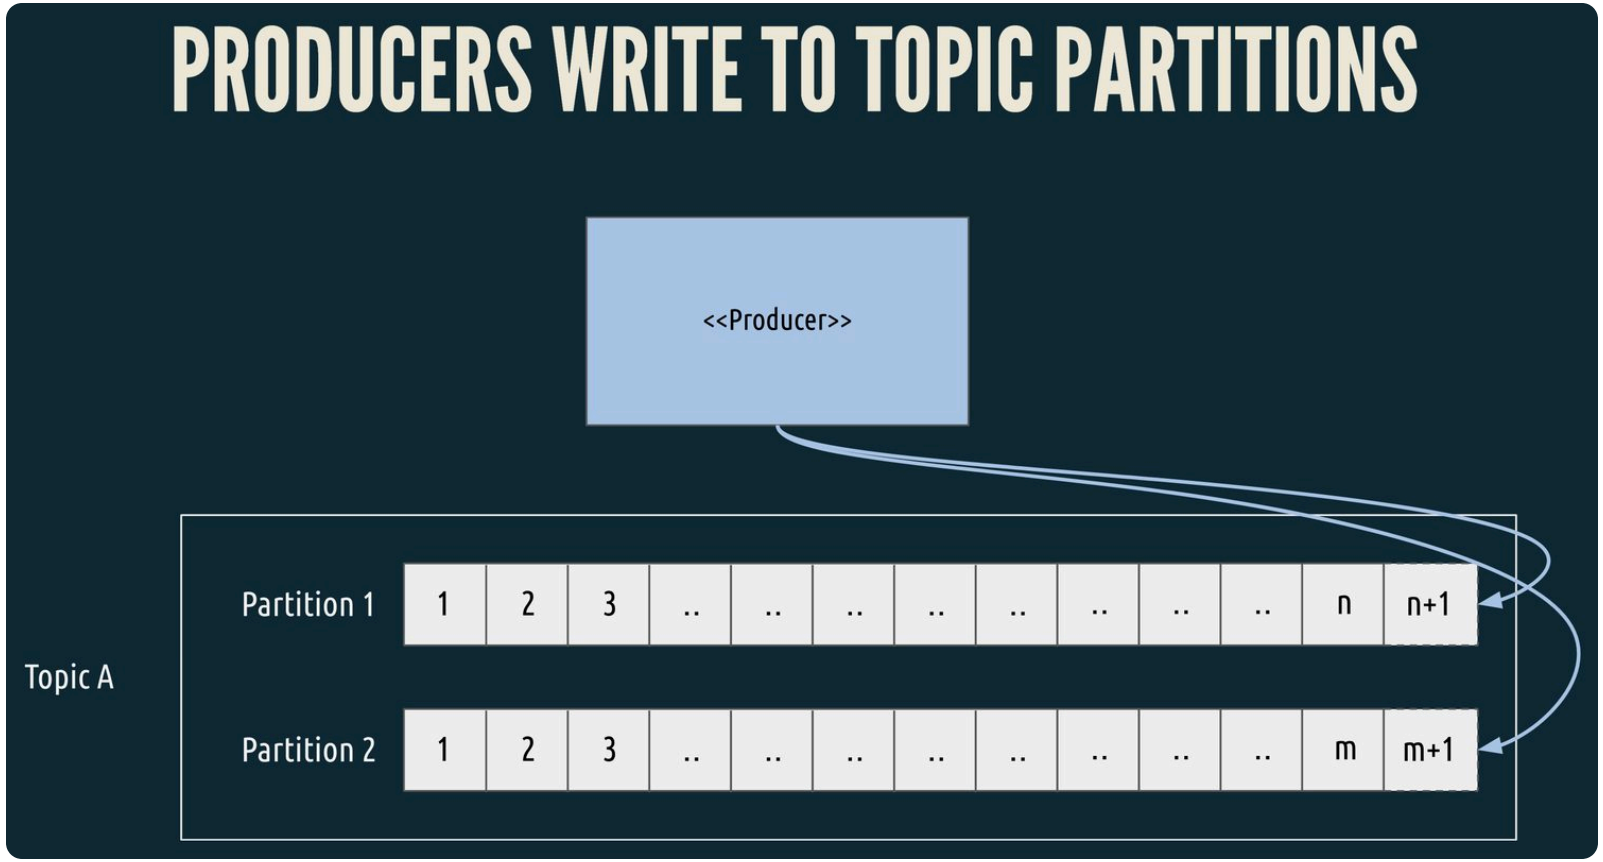
It's a log: an ordered, persistent sequence of events, **append-only**

Order Within a Partition



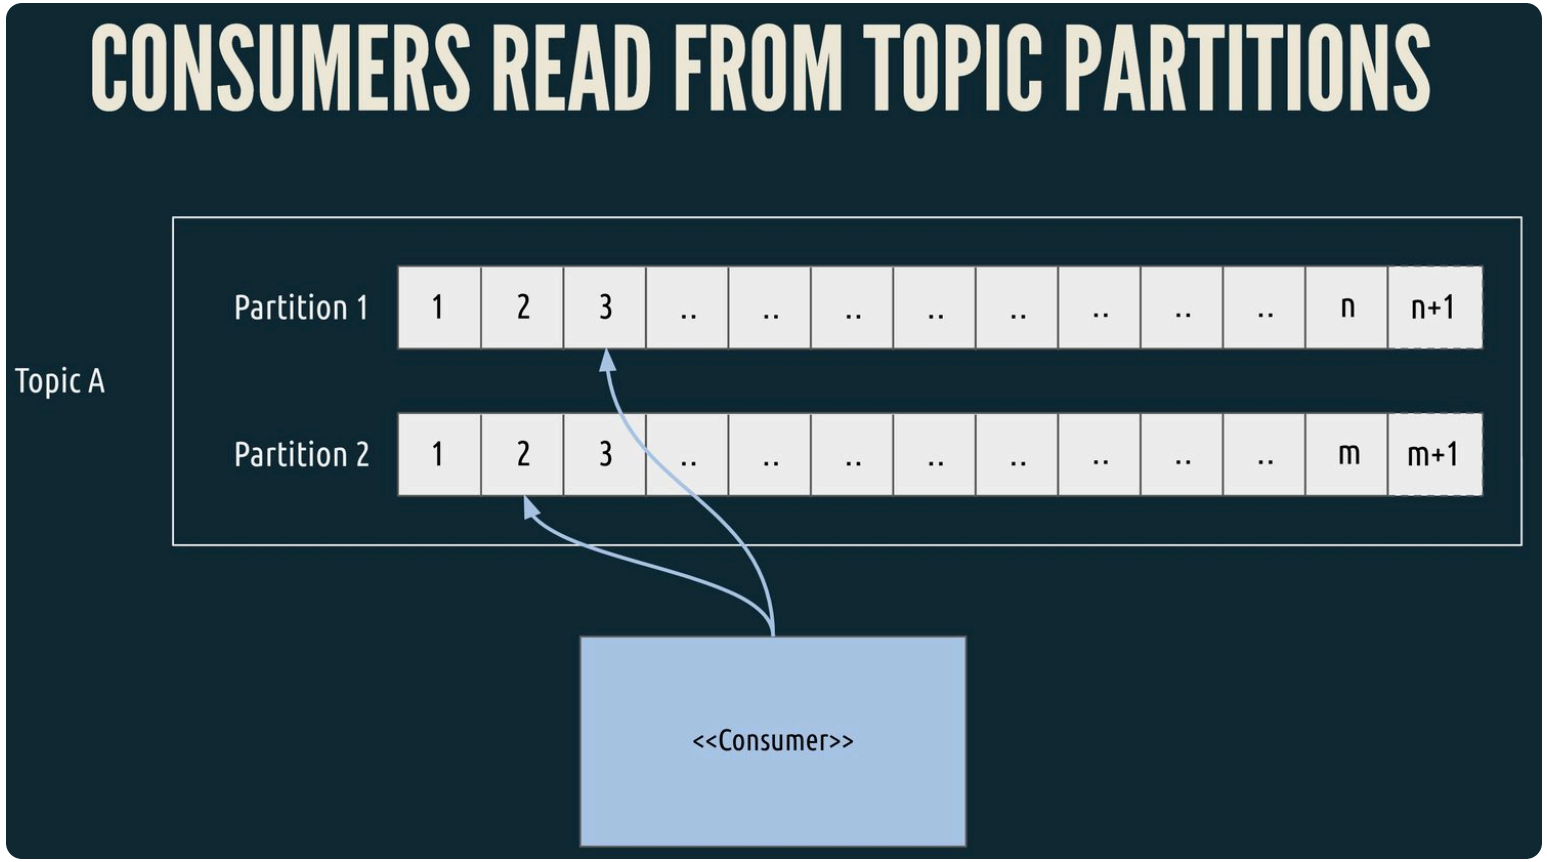
Topics are split into partitions: **order holds within one partition**, not across

Writing to Partitions



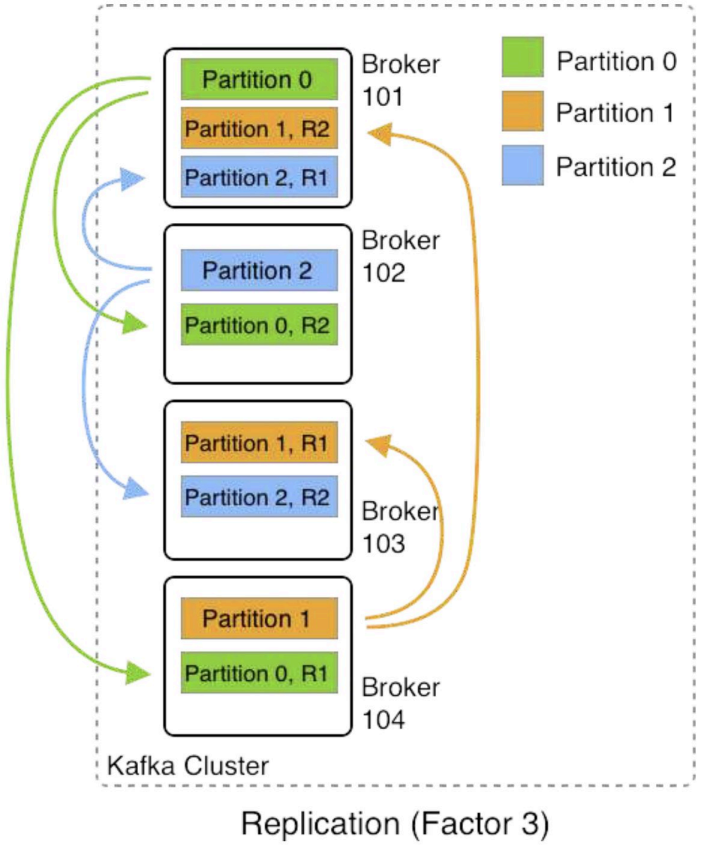
Producers write to topic partitions, **at the end**

Reading Without Deleting



Consumers read from topic partitions; **reading removes nothing**

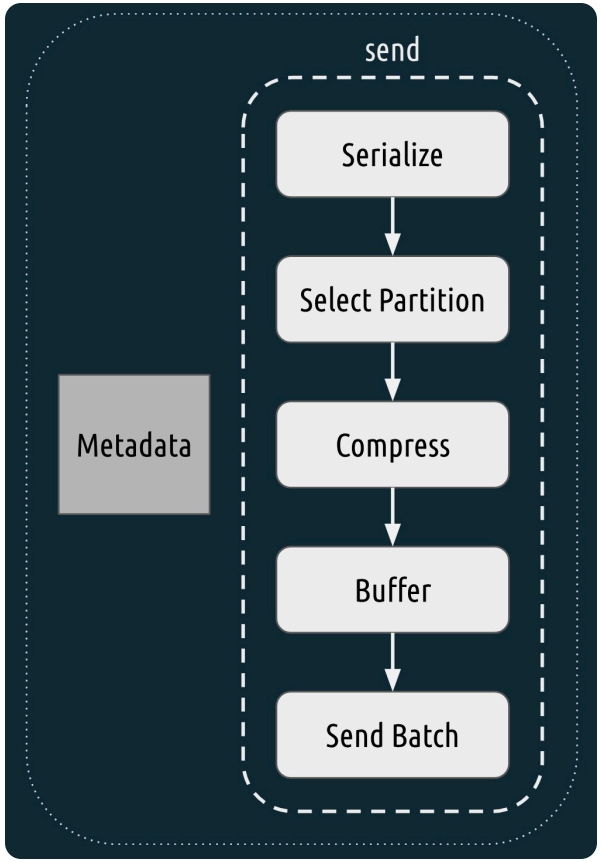
Broker Replication: Fault Tolerance



Replication (Factor 3)

Here each partition has a **leader and two followers** on different brokers

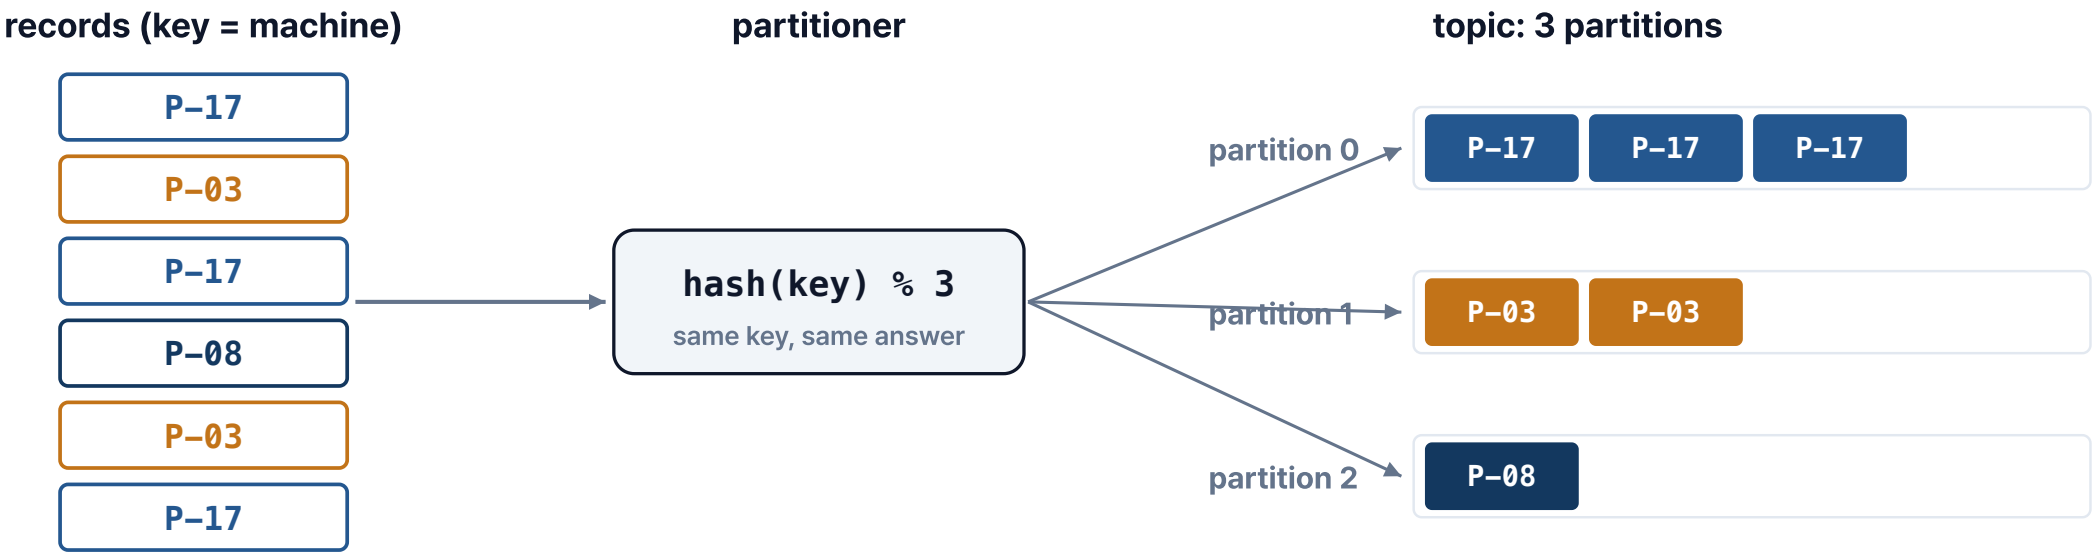
Producer



Serialize, pick a partition, compress, buffer, send a batch

Producer – Partitioning Strategy

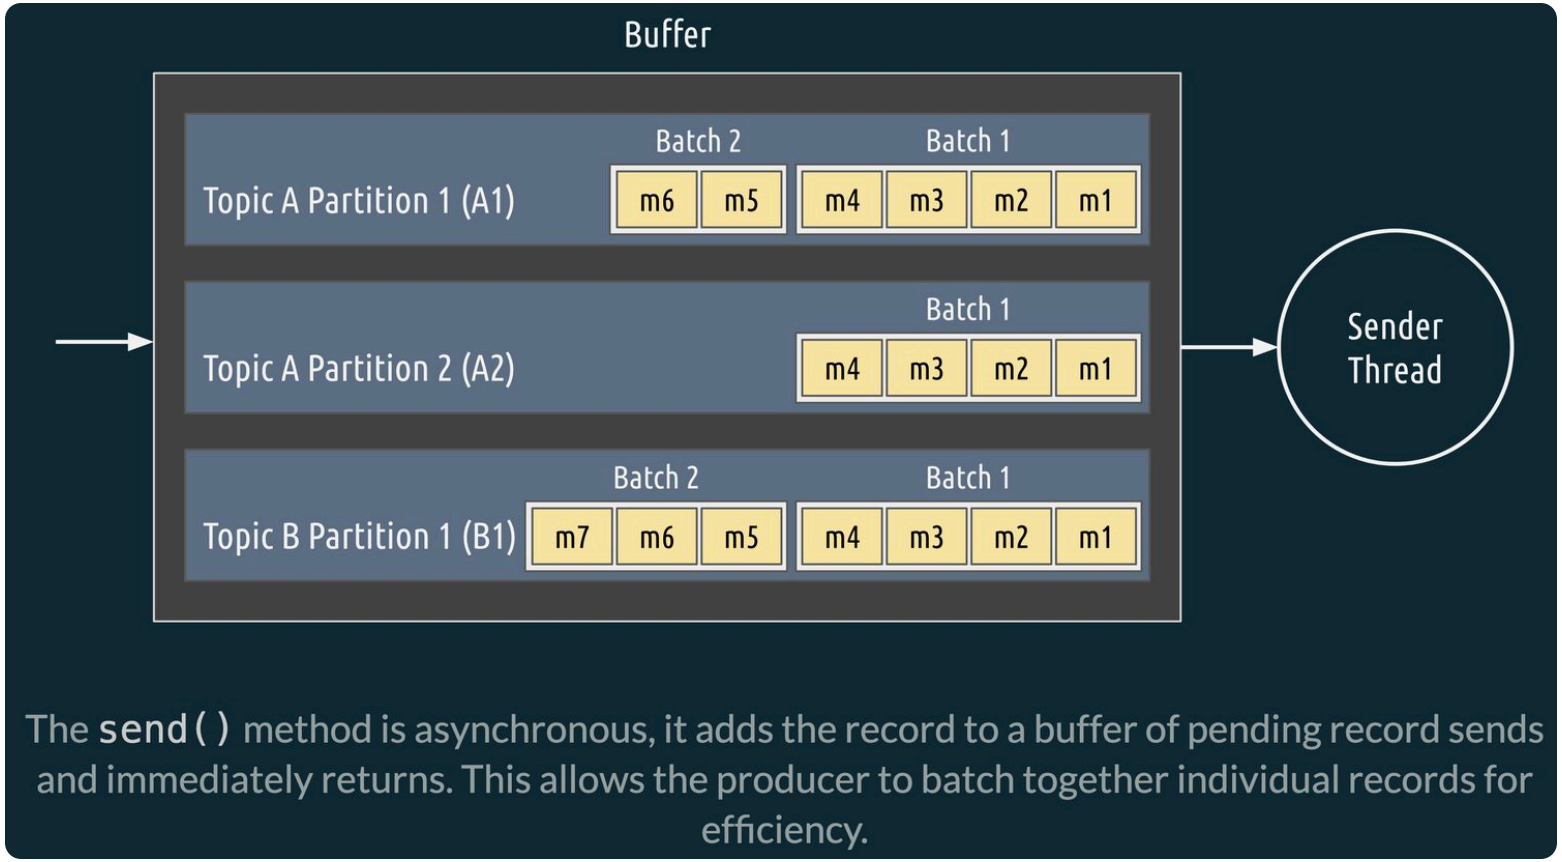
The record key decides the partition, so one key keeps its order



All P-17 events sit in one partition, **in the order they were sent**; no key → spread evenly

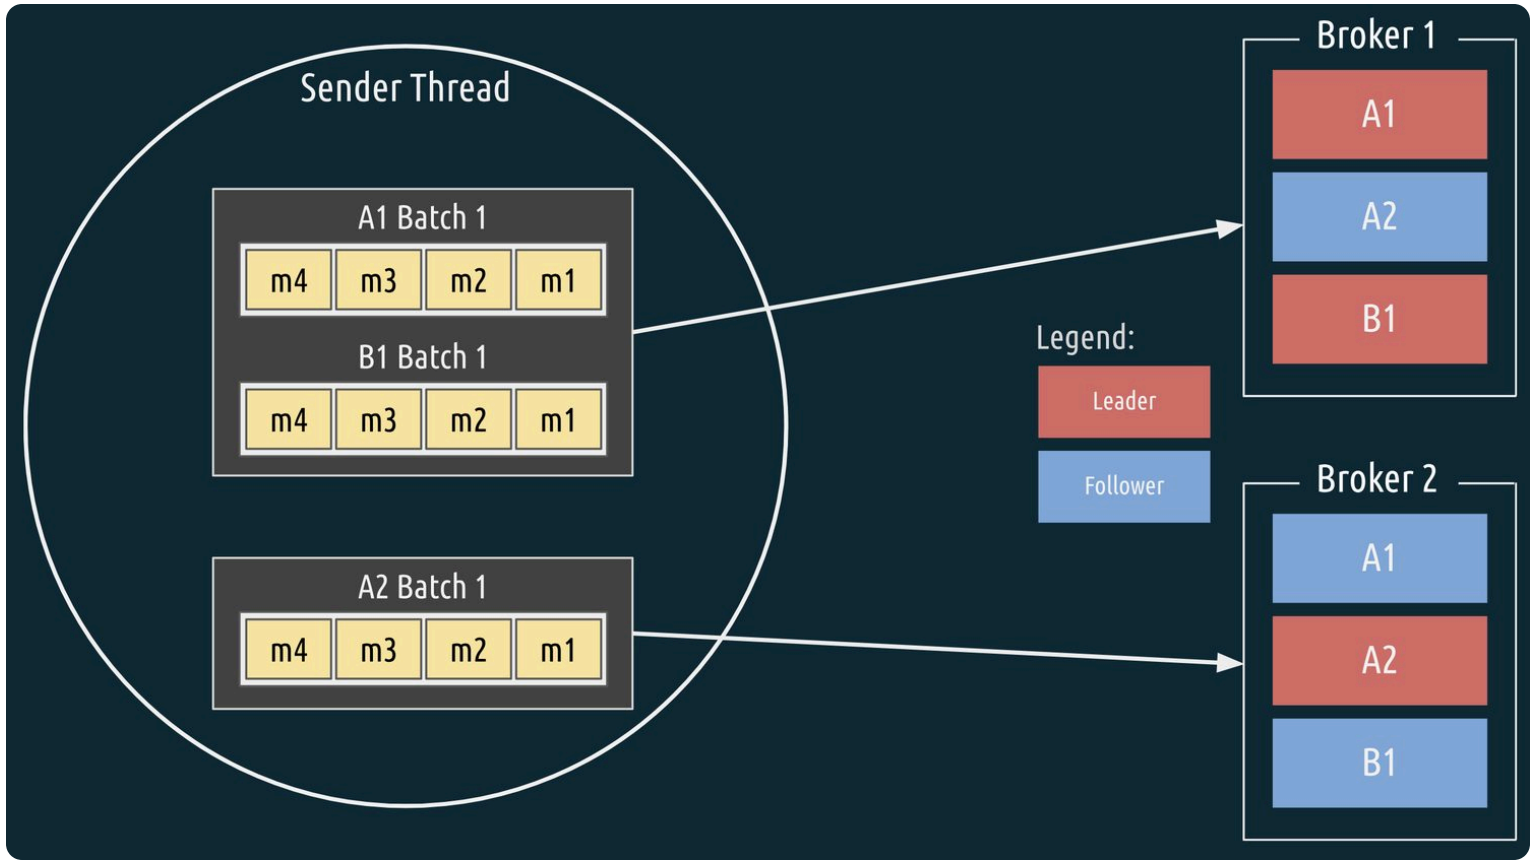
Same key, same partition: **order holds per key**

Producer – Accumulate Messages



send() returns at once; records **accumulate per partition** in the buffer

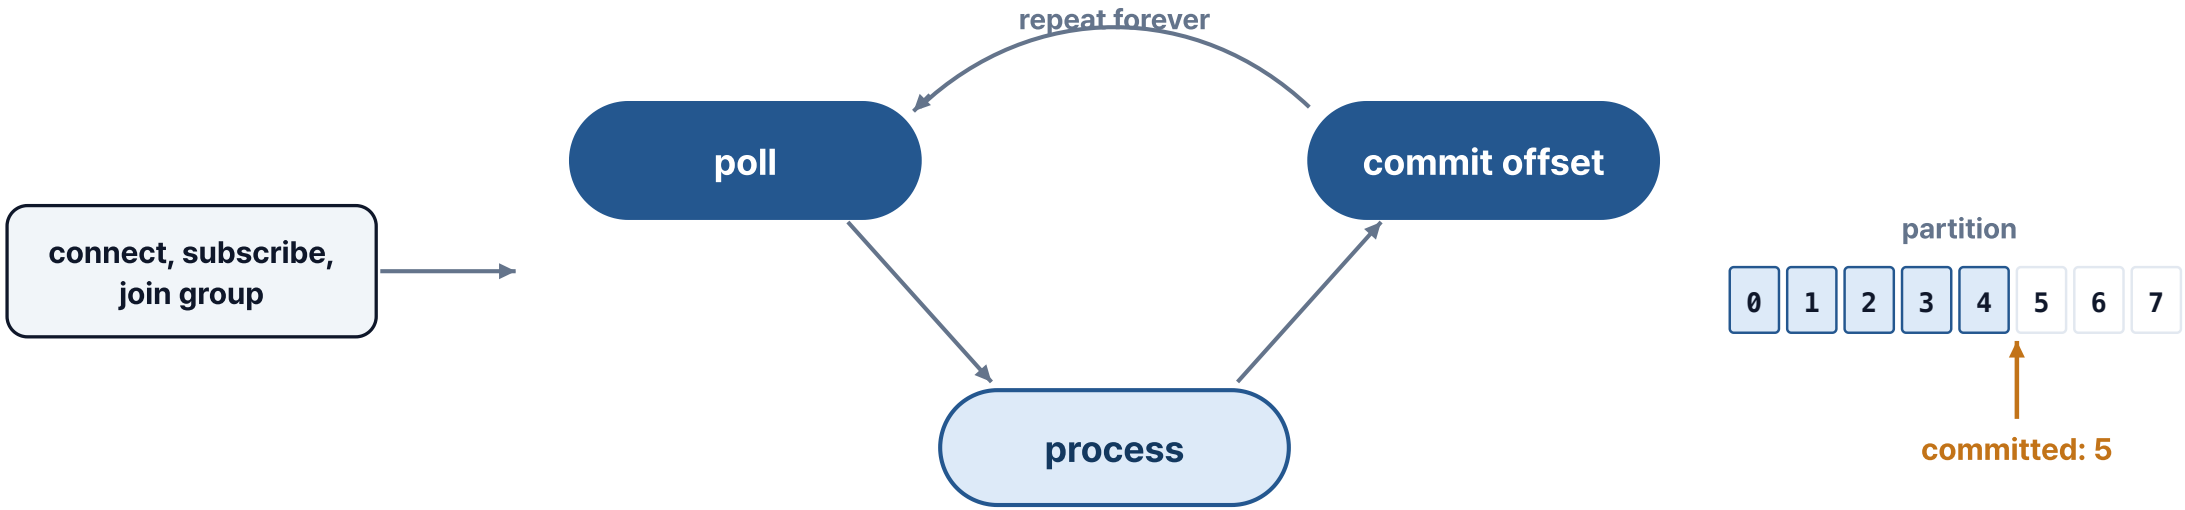
Producer – Send Message Batches



Batches go to each partition's **leader**; one request per broker

Consumer

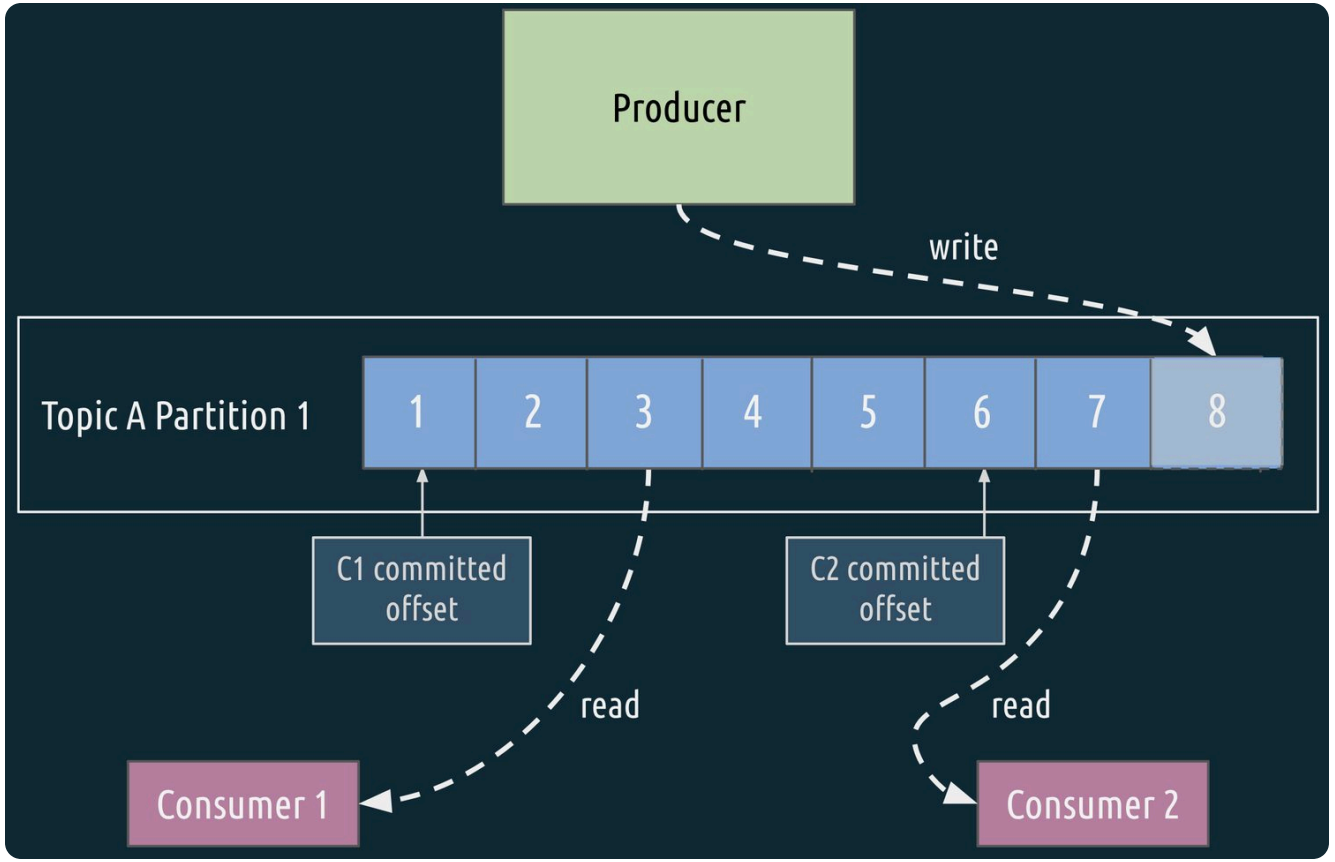
Join once, then poll, process and commit in a loop



After a crash the consumer resumes at the **committed offset**

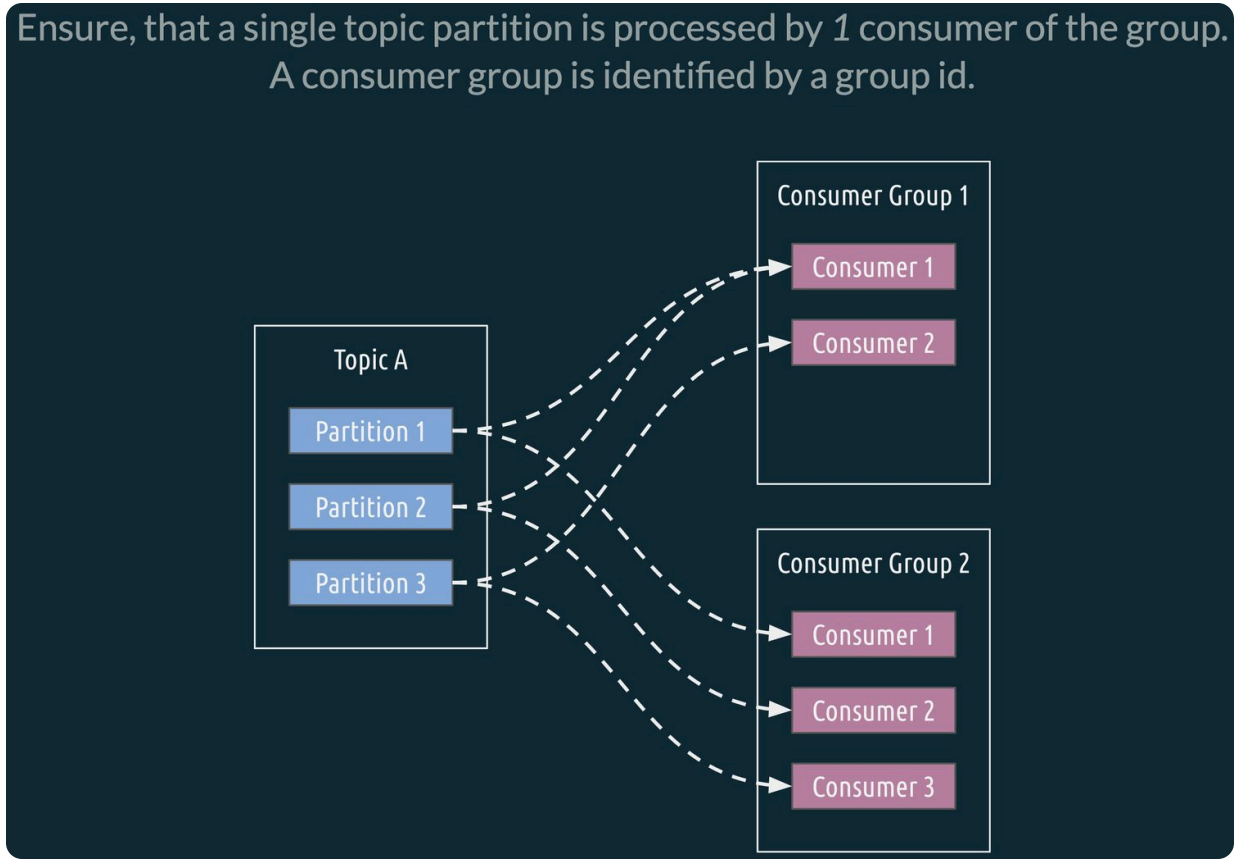
A restart resumes at the **last committed offset**

Consumer – Offsets



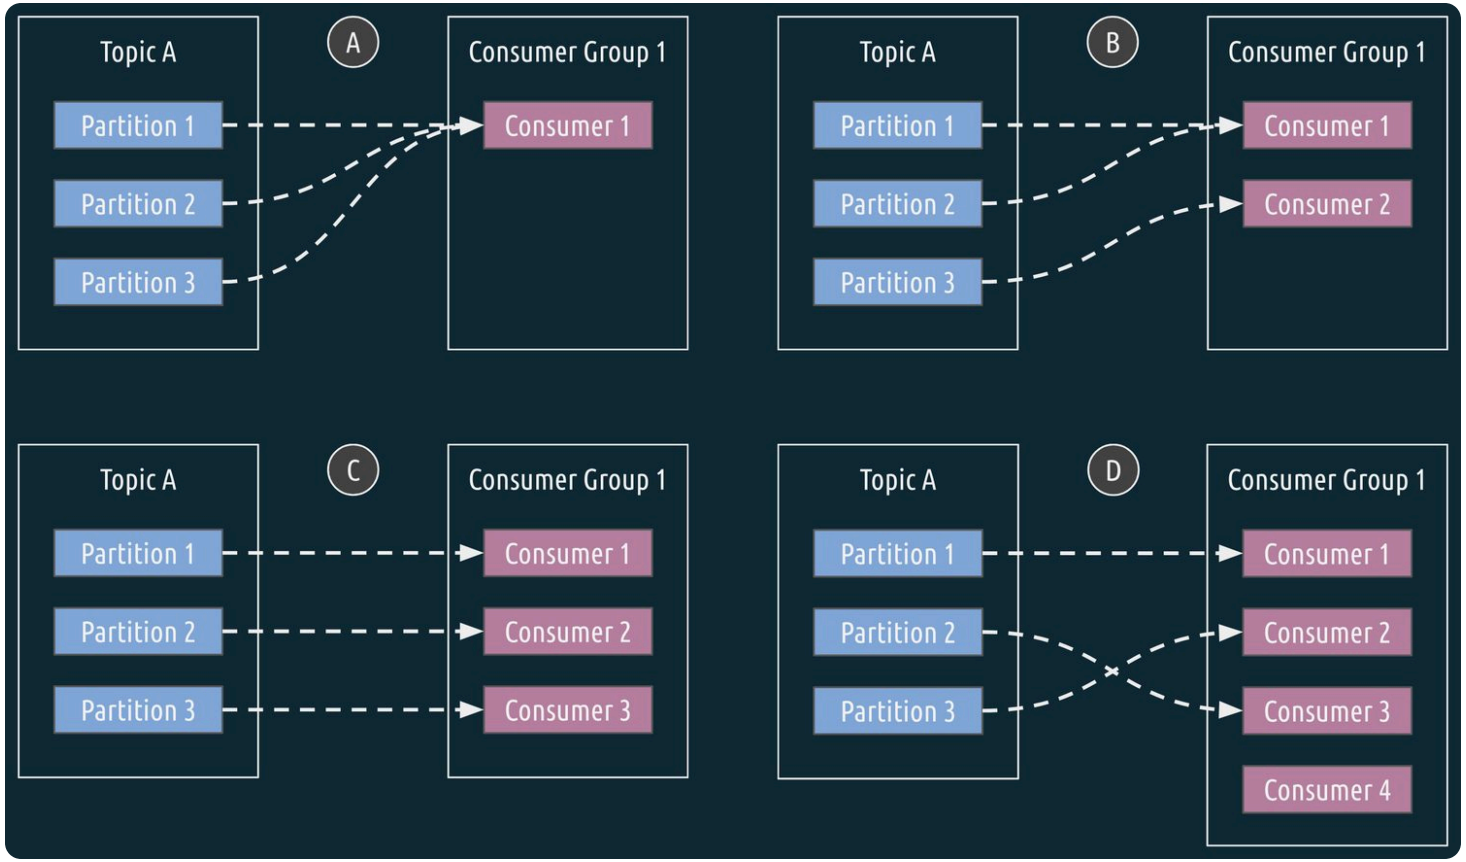
An offset is a position in the log; each consumer keeps **its own**

Consumer – Groups



Within a group a partition has **one reader**; different groups each read everything

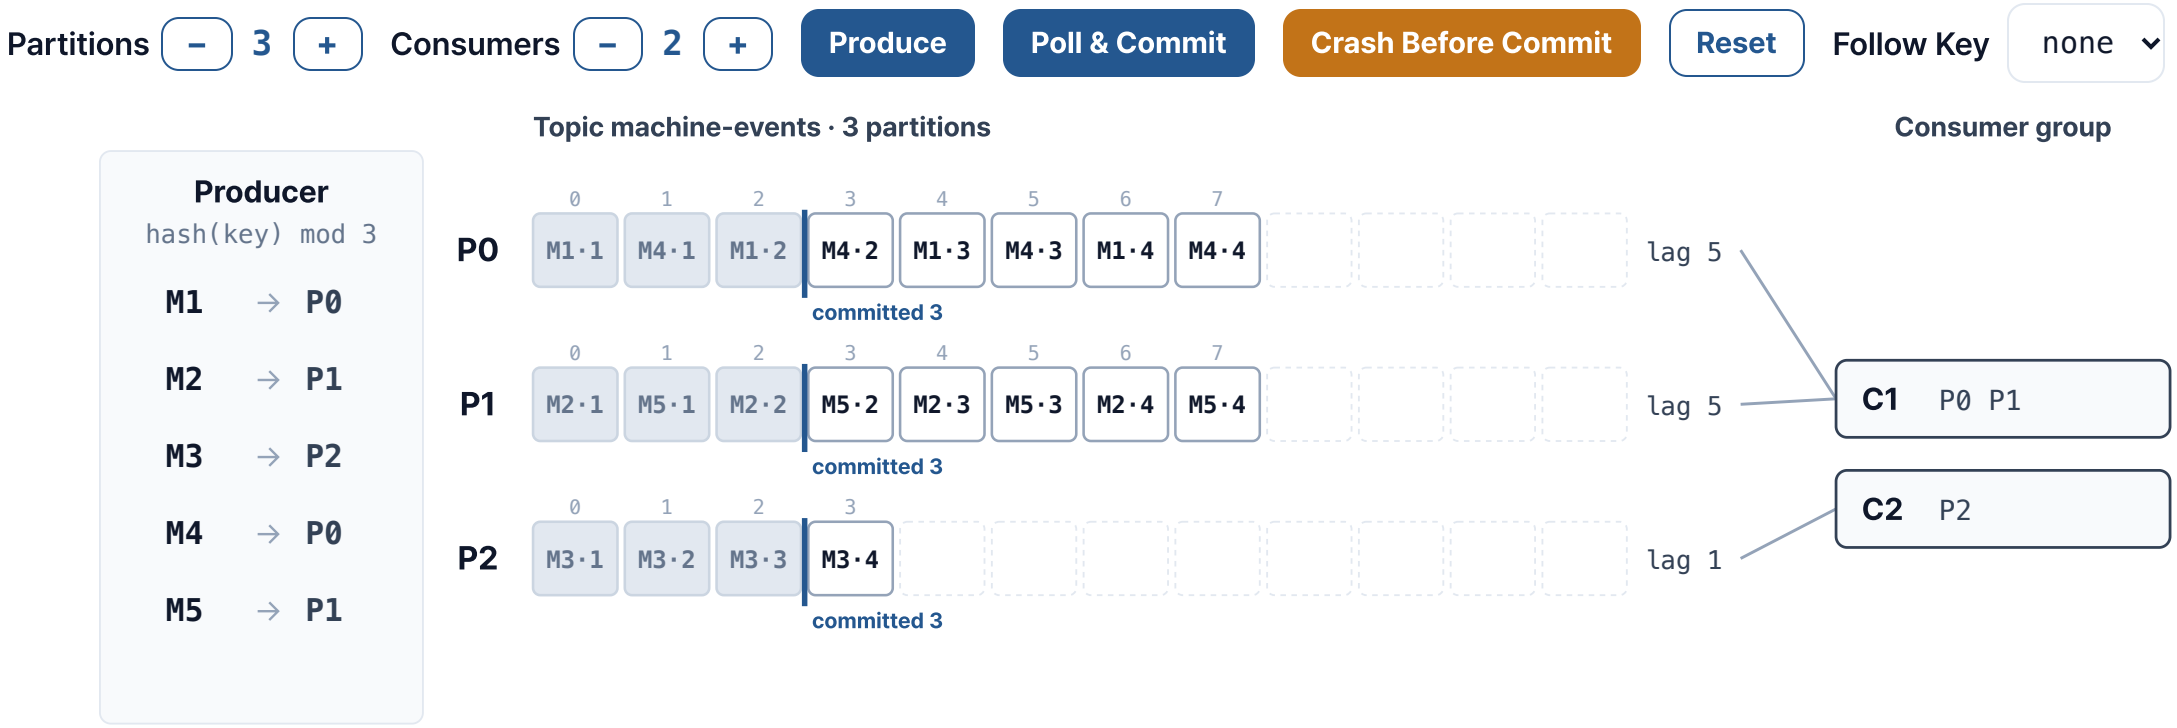
Consumer – Group Rebalancing



Partitions are reassigned as consumers come and go; **partitions cap the parallelism**

Demo: Partitions and Consumer Groups

Five machines send keyed events; change the group, then crash a consumer before it commits



Processed 9 · distinct 9 · **duplicates 0** · total lag 11

Start: 3 partitions, 2 consumers. C1 owns P0 and P1, C2 owns P2. Gray cells are processed and committed

Answer: Partitions and Consumer Groups

Explain the key order, parallelism limit, and replay after a crash

Change	Observation	Reason
Produce for M1	Same partition; increasing offsets	Key mapping is fixed
3 partitions; 4 consumers	One consumer is idle	One owner per partition in this group
Crash before commit; poll again	Previously processed records repeat	Restart from the committed offset

The stored log and the committed offset determine what is read

Caveat: At-Least-Once Duplicates

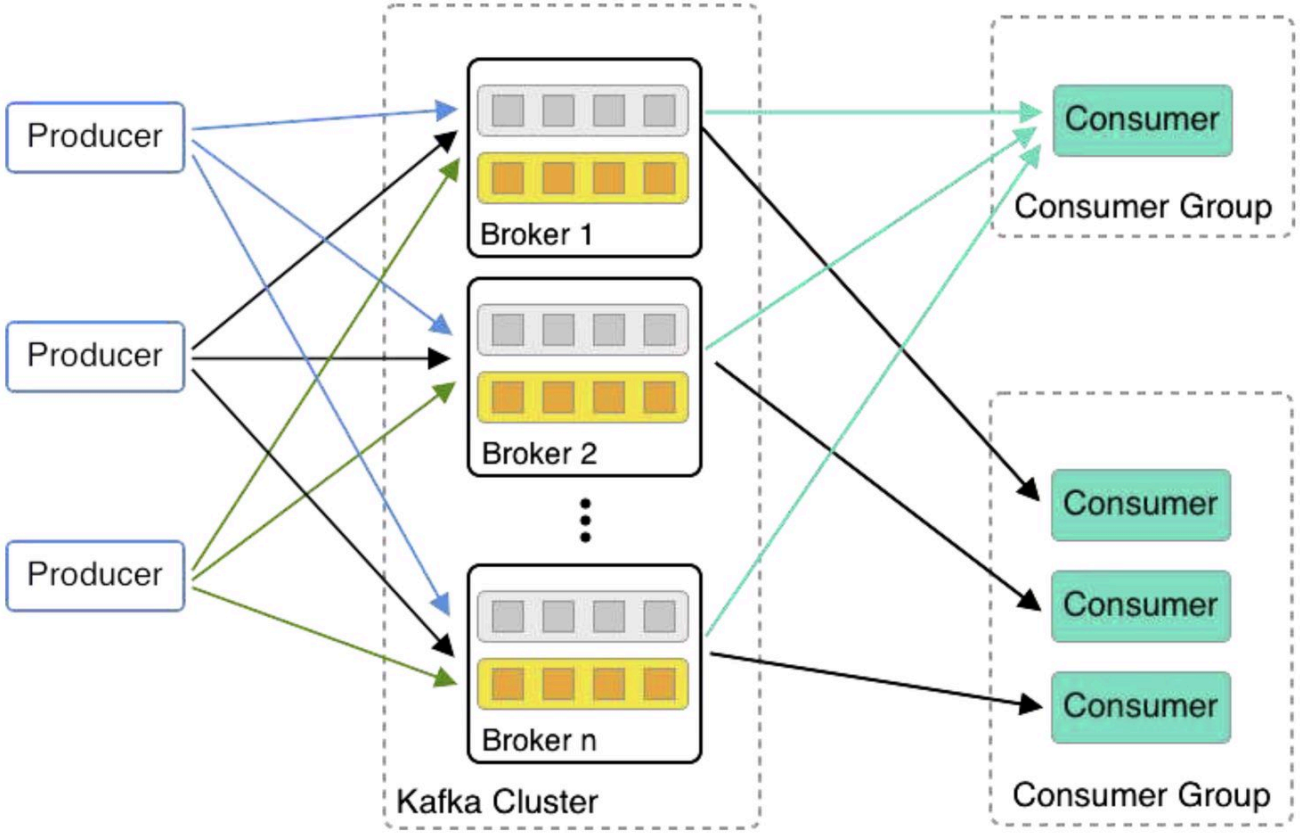
Process, then crash before the commit: the restart reads the same records again



Remedy: deduplicate by event ID, recorded in the same transaction as the side effect

The commit never happened, so offsets **4 to 6 are processed twice**

Conclusion: Scalable Data Pipeline

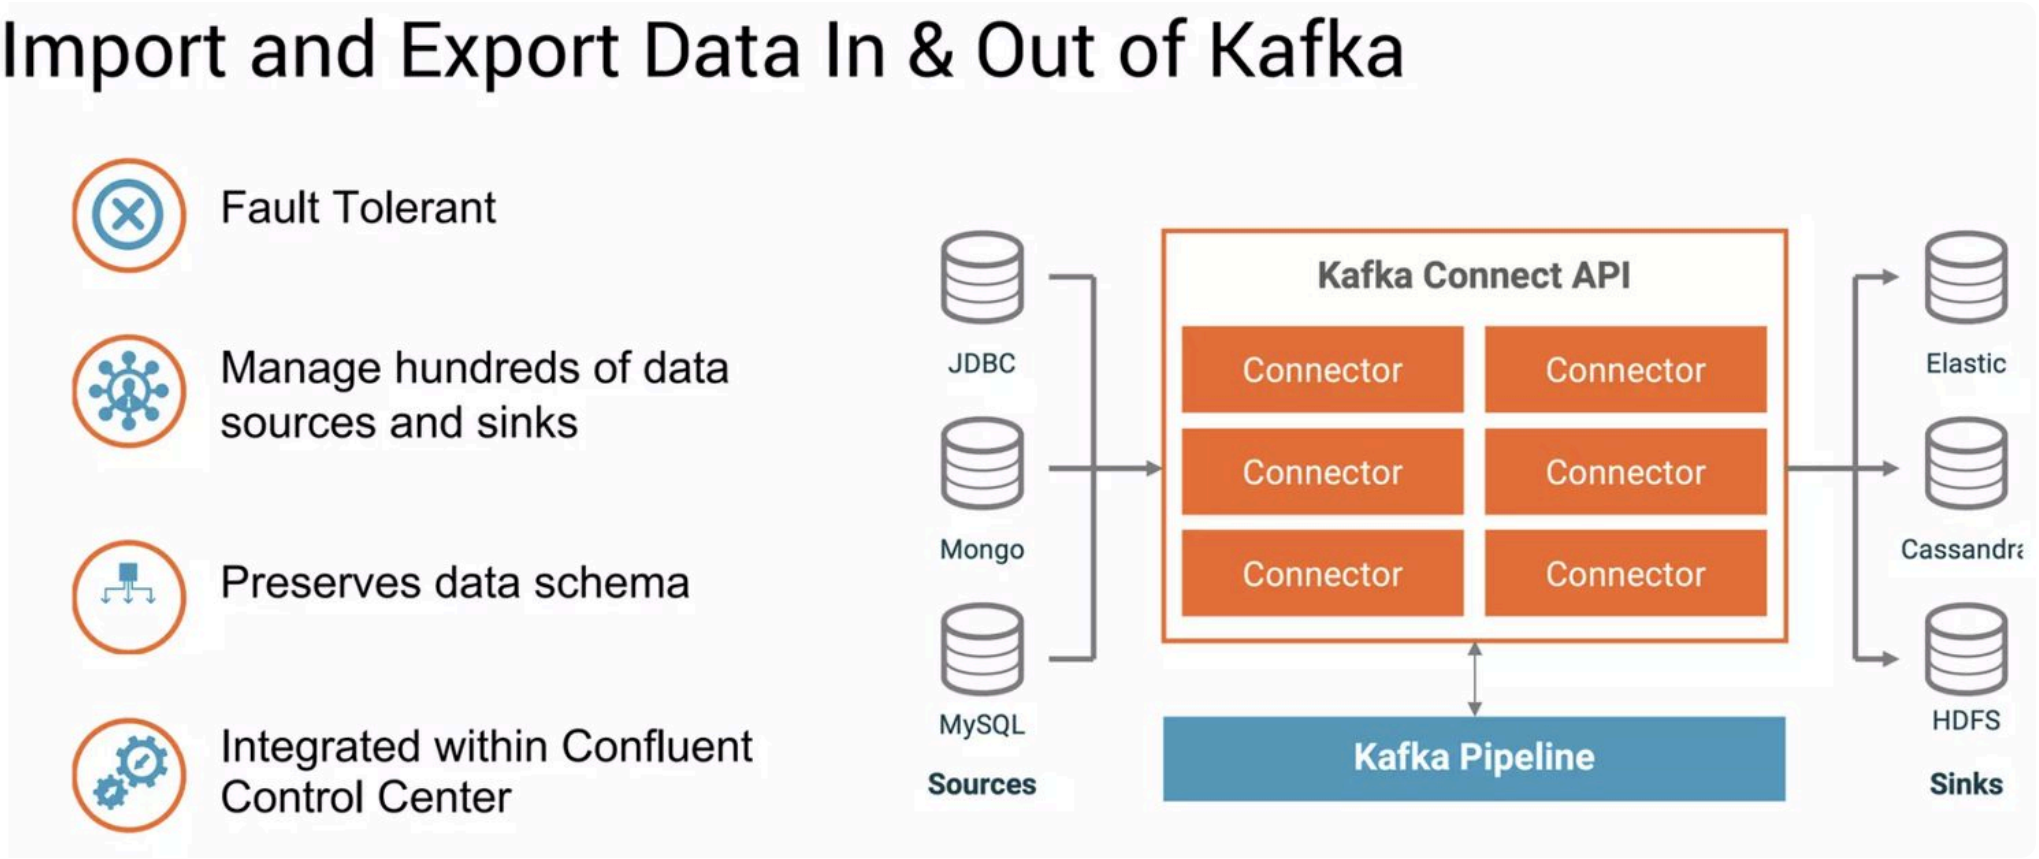


Add producers, brokers or consumers independently: **that is the point**

03 · Processing Event Streams

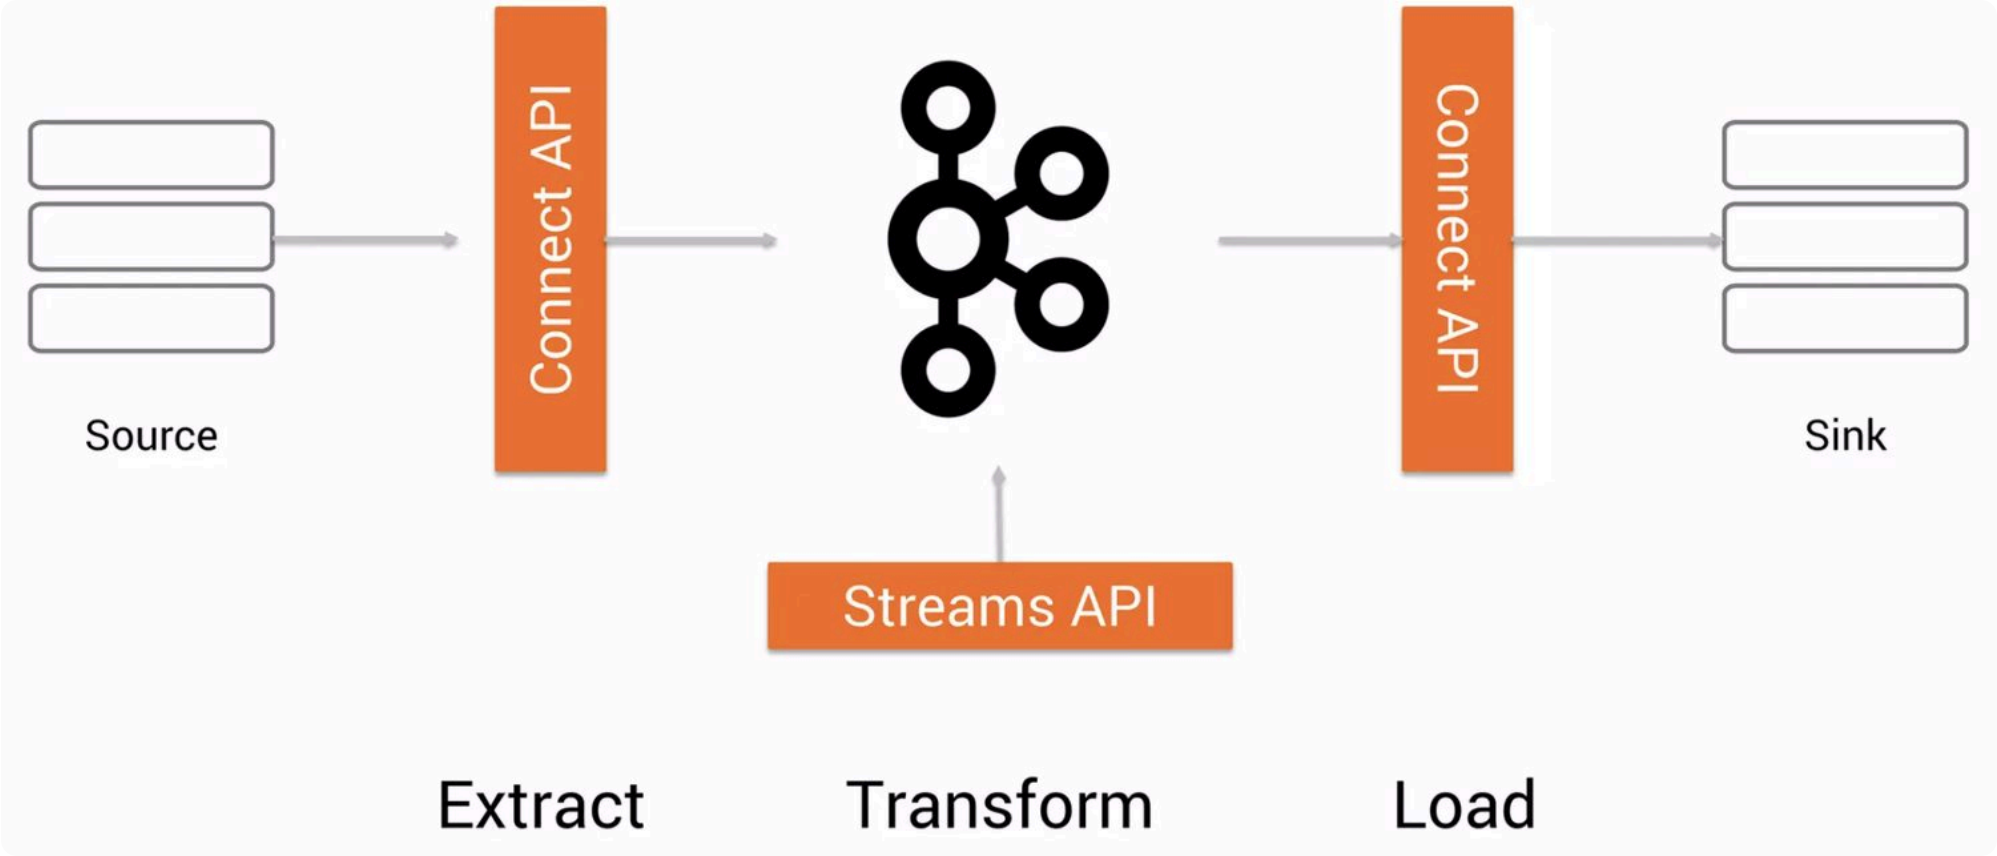
01	From Batch to Live Events	handle events as they arrive
02	Moving Events with Kafka	store events and deliver them reliably
03	**Processing Event Streams**	connect sources and process live data

Kafka Connect Sources and Sinks



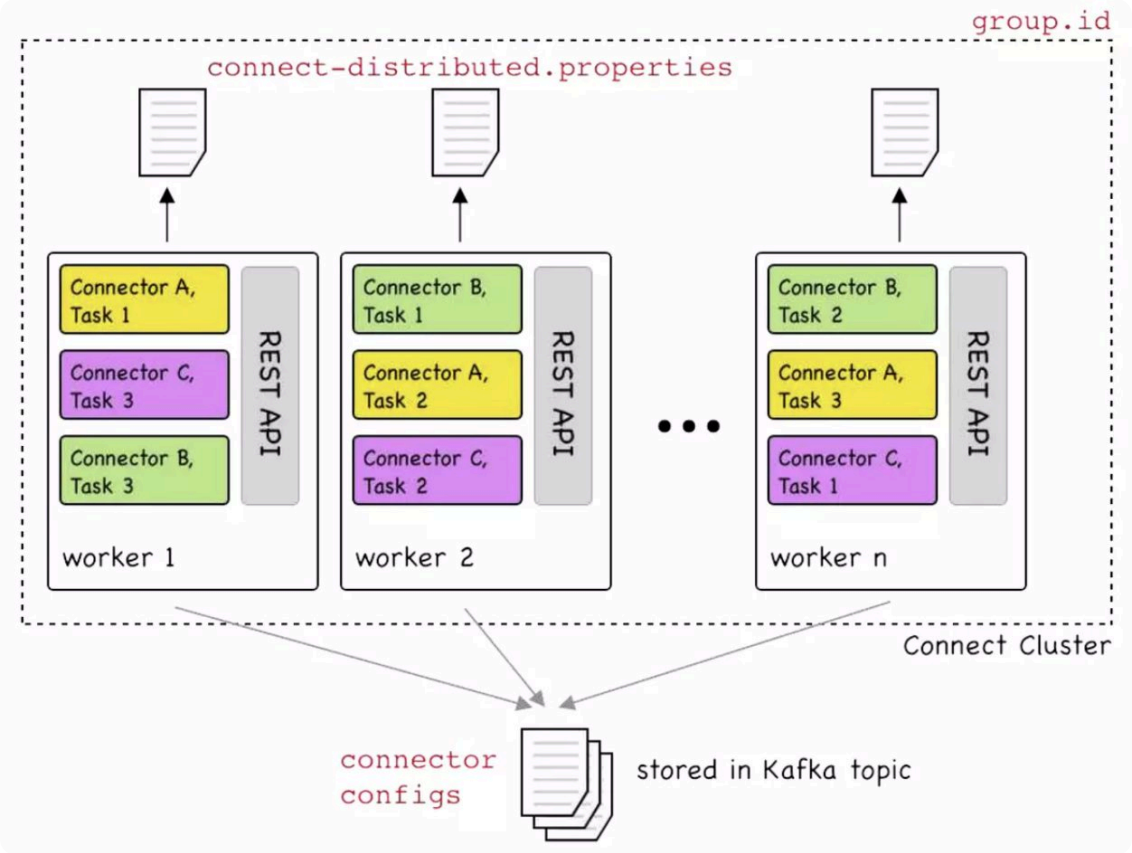
Source connectors in, sink connectors out: **configuration, not code**

A Continuous ETL Flow



Extract (source connector), transform (Streams), load (sink connector)

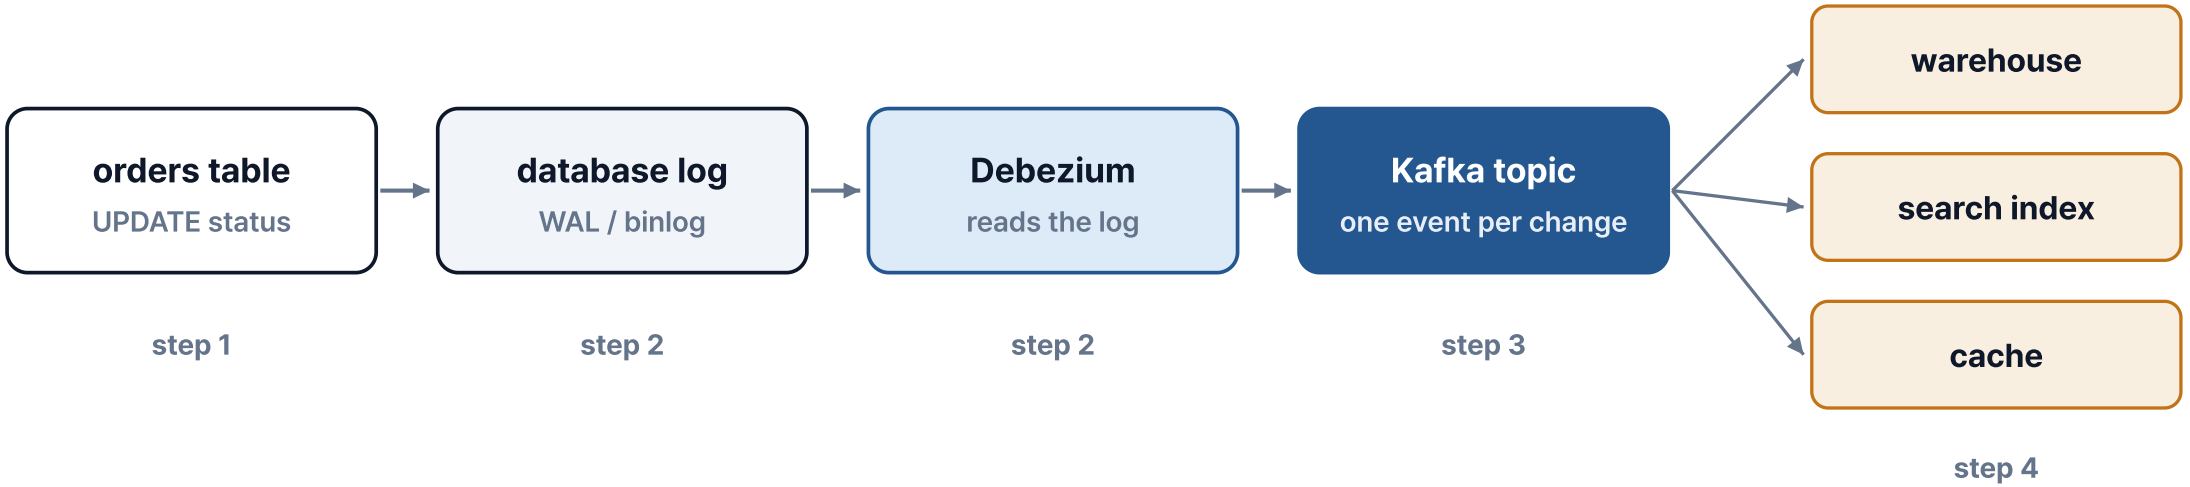
Kafka Connect Workers



Connector tasks spread over workers; configs live **in a Kafka topic**

Change Data Capture

Read the database log and publish every row change as an event



Only changed rows move, in commit order, seconds after they happen

Every copy stays in sync by **replaying the changes**, not re-copying tables

Kafka Streams as a Library

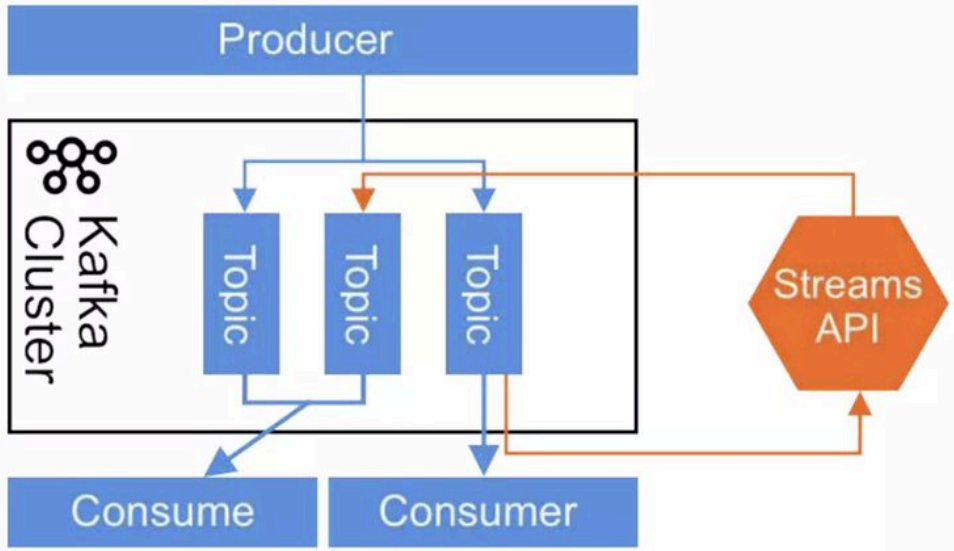
Transform Data with Real-Time Applications

Overview

- Write standard Java applications
- No separate processing cluster required
- Exactly-once processing semantics
- Elastic, highly scalable, fault-tolerant
- Fully integrated with Kafka security

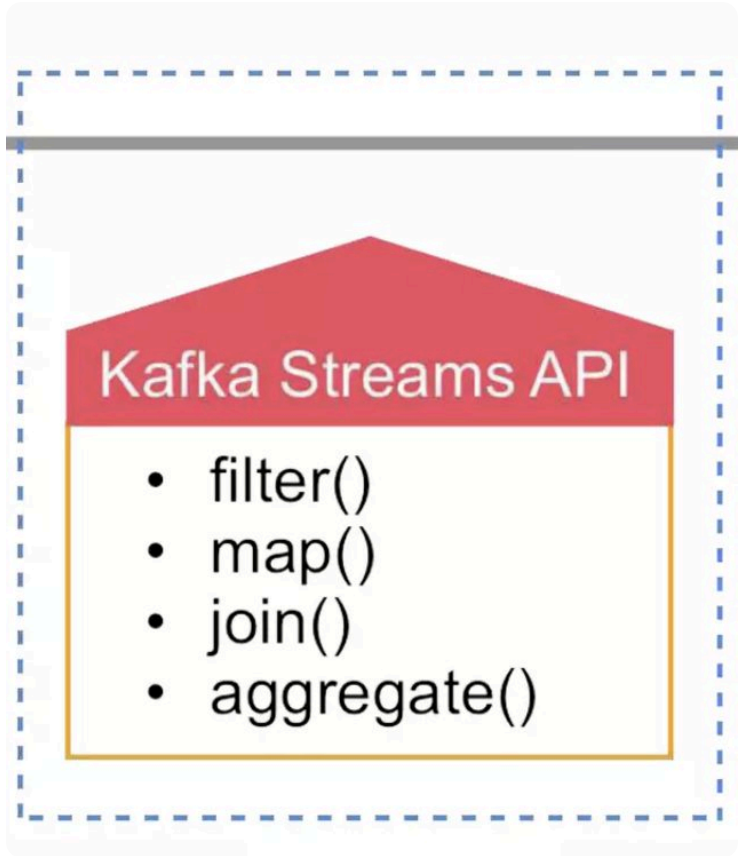
Example Use Cases

- Microservices
- Continuous queries
- Continuous transformations



A library, not a cluster: your app reads a topic, transforms, writes a topic

Continuous Transformations



filter, map, join, aggregate, **on an unbounded stream**

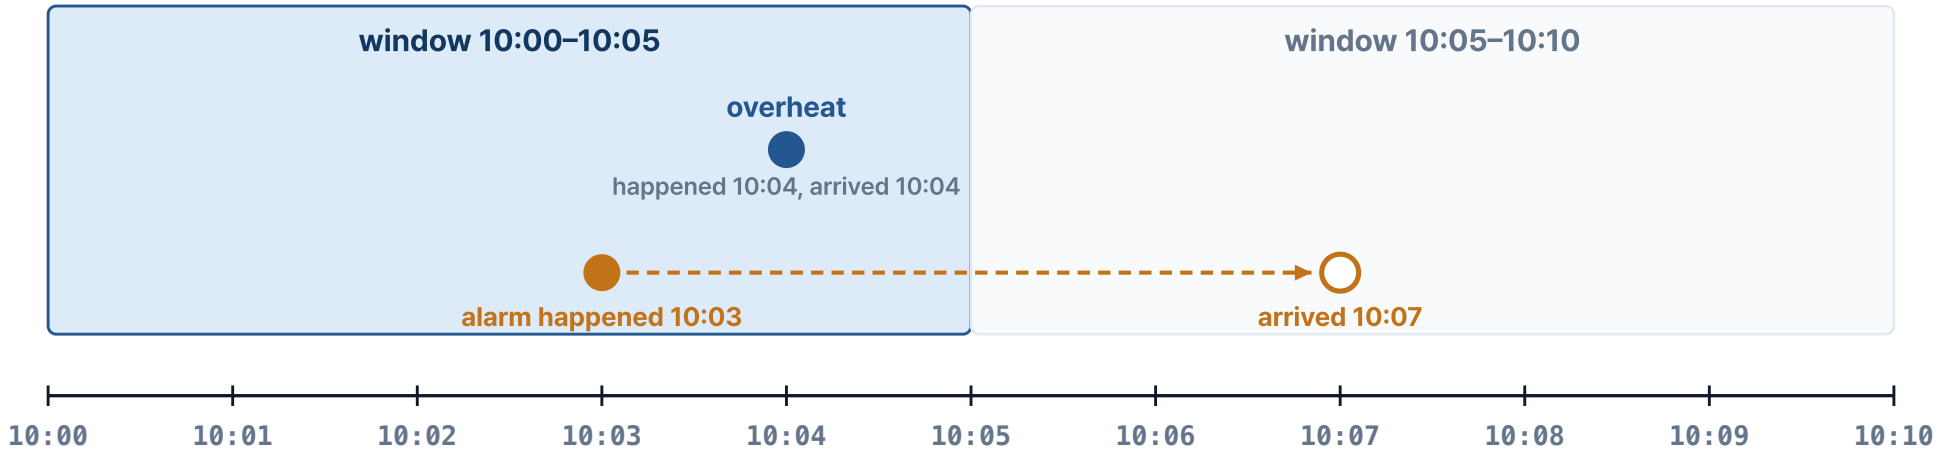
Confluent KSQL

Develop real-time stream processing apps writing only SQL!
No Java, Python, or other boilerplate to wrap around it.

```
CREATE STREAM vip_actions AS
  SELECT   userid, page, action
  FROM     clickstream c
  LEFT JOIN users u ON c.userid = u.user_id
  WHERE    u.level = 'Platinum';
```

Event Time vs Arrival Time

An event can arrive after the time window in which it happened



By **event time** the alarm belongs to the first window; by arrival time it would not

The alarm belongs to the earlier window by **event time**

Watermarks and Late Events

A later close improves completeness but keeps window state longer



Waiting longer gives a **more complete** answer, **later**, and keeps more state

Latency and completeness move together when events arrive late

Stream Processing Engines

Engine	Runs as	Good at
Kafka Streams	a library inside your app	small transforms next to Kafka
Apache Flink	its own cluster	large stateful jobs, event-time windows, exactly-once
Spark Structured Streaming	Spark, in small batches	teams that already run Spark

Kafka moves the events; **Kafka Streams, Flink or Spark** compute on them

Example – Create Topic

Before you can write your first events, you must create a topic. Open another terminal session and run:

```
$ bin/kafka-topics.sh --create --topic quickstart-events --bootstrap-server
localhost:9092
```

All of Kafka's command line tools have additional options. Run `kafka-topics.sh` without arguments to display usage information, or describe the new topic:

```
$ bin/kafka-topics.sh --describe --topic quickstart-events --bootstrap-server
localhost:9092
Topic: quickstart-events   TopicId: NPmZHyhbR9y00wWglMH2sg   PartitionCount: 1
ReplicationFactor: 1
  Topic: quickstart-events   Partition: 0   Leader: 0   Replicas: 0   Isr: 0
```

Example – Process Events with Kafka Streams

To give you a first taste, here's how one would implement the popular `WordCount` algorithm:

```
KStream<String, String> textLines = builder.stream("quickstart-events");

KTable<String, Long> wordCounts = textLines
            .flatMapValues(line -> Arrays.asList(line.toLowerCase().split(" ")))
            .groupBy((keyIgnored, word) -> word)
            .count();

wordCounts.toStream().to("output-topic", Produced.with(Serdes.String(),
Serdes.Long()));
```

The same WordCount as Lecture 11's MapReduce and Spark, now **updated with every new line**: a KTable is a count that never finishes.

Takeaway: Big Data III

1. Kafka stores ongoing events in **partitioned logs**; order holds within a partition
2. Consumer groups scale reading, but a crash before an **offset commit** can replay events
3. Connect moves events; a stream processor computes by **event time** and must handle late arrivals